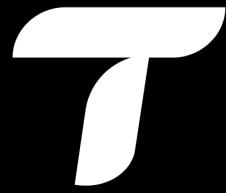




1000X WORLD TOUR

13 CITIES | 6 DAYS





TigerBeetle

How I Downgraded a Safety Bug to an
Availability bug... using Assertions!

December 4th, 2025

Consensus Protocols

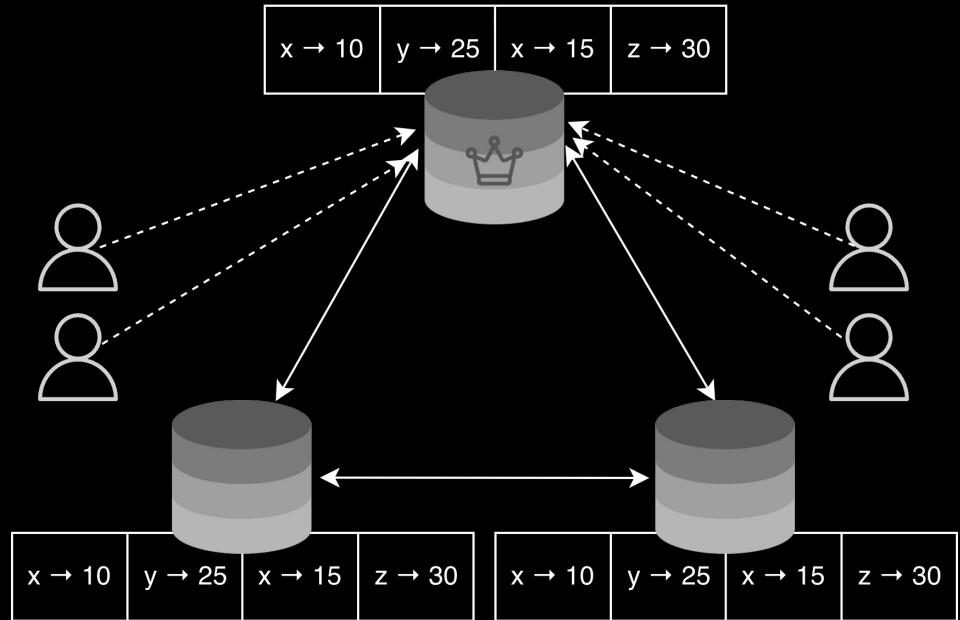
Agreement on the state of the system, even under faults

What do these protocols guarantee?

Fault Tolerance

Ordering

Strict Serializability



Building Consensus Protocols

Notoriously hard to get right

How we test consensus at τ

Deterministic Simulation Testing

- Nondeterministic, physical interactions stubbed
- Time sped up
- Finds bugs using *assertions*



Assertions

Validate system invariants at runtime

Assertions across layers

Assertions in Prod

- Enforce *replica*-level invariants

Assertions in DST

- Check *cluster*-level invariants

```
fn on_request_start_view(
    self: *Replica,
    message: *const Message.RequestStartView,
) void {
    assert(message.header.command == .request_start_view);
    if (self.ignore_repair_message(message.base_const())) return;

    assert(self.status == .normal);
    assert(self.view == self.log_view);
    assert(message.header.view == self.view);
    assert(message.header.replica != self.replica);
    assert(self.primary());

    const start_view_message = self.create_start_view_message(message.header.nonce);
    defer self.message_bus.unref(start_view_message);

    assert(start_view_message.header.command == .start_view);
    assert(start_view_message.references == 1);
    assert(start_view_message.header.view == self.view);
    assert(start_view_message.header.op == self.op);
    assert(start_view_message.header.commit_max == self.commit_max);
    assert(start_view_message.header.nonce == message.header.nonce);
    self.send_message_to_replica(message.header.replica, start_view_message);
}
```

those systems. The systems were considered successful, yet bugs and operational problems persisted. To mitigate the problems, the systems used well-proven methods—pervasive contract assertions enabled in production—to detect symptoms of bugs, and mechanisms (such as “recovery-oriented computing”²⁰) to attempt to minimize the impact when bugs are triggered.

How Amazon Web Services Uses Formal Methods” – Newcombe et al., CACM ’15

JEPSEN

TigerBeetle 0.16.11

Kyle Kingsbury
2025-06-06

TigerBeetle is a distributed OLTP database oriented towards financial transactions. We tested TigerBeetle 0.16.11 through 0.16.30. We discovered seven client and server crashes, including a segfault on client close and several panics during server upgrades. Single-node failures could cause significantly elevated latencies for the duration of the fault, and requests were intentionally retried forever, which complicates error handling. We found only two safety issues: missing results for queries with multiple predicates, and a minor issue with a debugging API returning incorrect timestamps. TigerBeetle offered exceptional resilience to disk corruption, including damage to every replica's files. However, it lacked a way to handle the total loss of a node's data. As of version 0.16.30, TigerBeetle appeared to meet its promise of Strong Serializability. As of 0.16.45, TigerBeetle had addressed every issue we found, with the exception of indefinite retries. TigerBeetle has written a companion blog post to this work. This report was funded by TigerBeetle, Inc., and conducted in accordance with the Jepsen ethics policy.

Root-cause study. About two thirds of software upgrade failures are caused by incompatible interaction between two software versions. The interaction occurs through either persistent data (60%) or network messages (40%), with the latter being a particular concern during rolling upgrades. Incompatible assumption about data syntax or semantics causes one version to fail to parse (about two thirds of the incompatibility) or handle (about one third of the incompatibility) storage files or network messages generated by another version. Our detailed study provides guidance on how to automatically detect incompatibilities and how to avoid cross-version incompatibilities during programming. More details are presented in section 4.

**Understanding and Detecting Software Upgrade
Failures in Distributed Systems - Zhang et al.,
SOSP 2021**

Primary

Backup

Upgrade

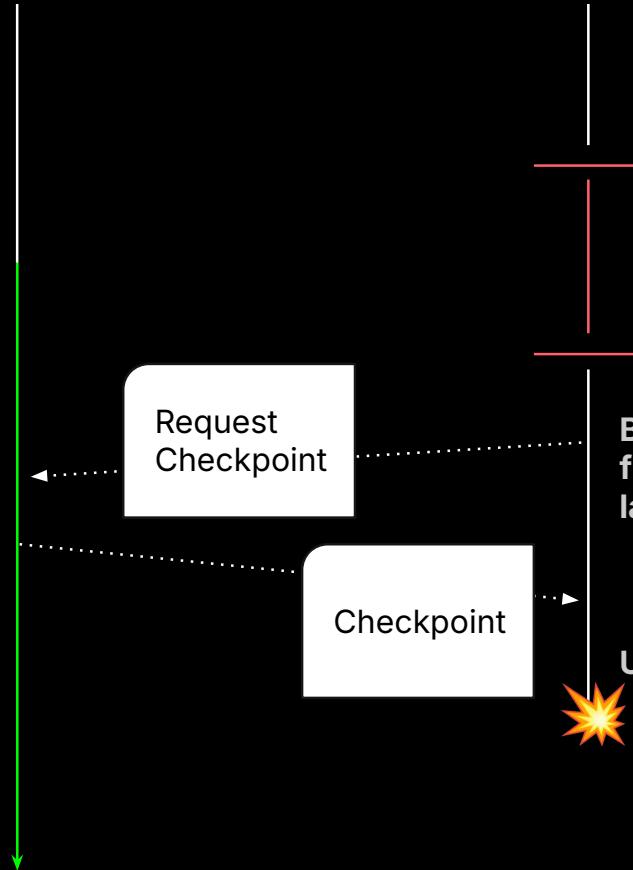


Request
Checkpoint

Backup
finds that its
lagging

Checkpoint

Upgrade



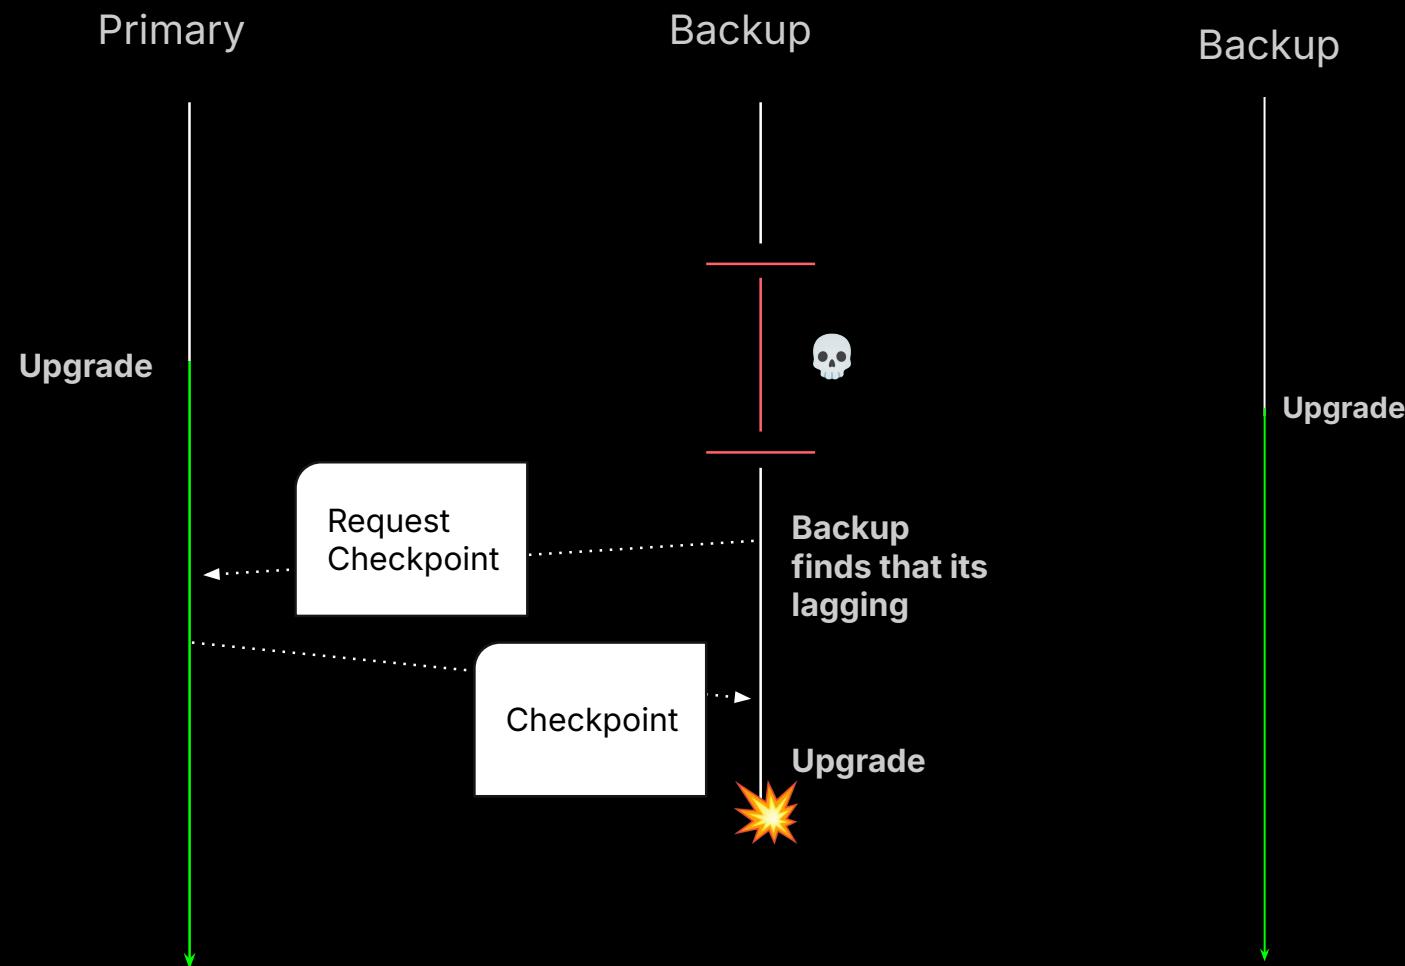
```
pub fn assert_free_set_consistent(self: *const Replica) void {
    assert(self.grid.free_set.opened);
    assert(self.state_machine.forest.manifest_log.opened);

    // Must be invoked either on startup, or after checkpoint completes.
    assert(!self.state_machine_opened or self.commit_stage == .checkpoint_superblock);

    var forest_tables_iterator = ForestTableIterator{};
    var tables_index_block_count: u64 = 0;
    var tables_value_block_count: u64 = 0;
    while (forest_tables_iterator.next(&self.state_machine.forest)) |table| {
        const block_value_count = switch (Forest.tree_id_cast(table.tree_id)) {
            inline else => |tree_id| self.state_machine.forest.tree_for_id_const(
                tree_id,
            ).block_value_count_max(),
        };
        tables_index_block_count += 1;
        tables_value_block_count += stdx.div_ceil(
            table.value_count,
            block_value_count,
        );
    }
}

assert((self.grid.free_set.count_acquired() - self.grid.free_set.count_released()) ==
       (tables_index_block_count + tables_value_block_count +
        self.state_machine.forest.manifest_log.log_block_checksums.count));
```

The assertion that caught the bug;
using free-set consistency to assert
storage determinism



Vulnerable Context. During upgrade, the system has to go through a no-service (full-stop upgrade) or partial-service (rolling upgrade) period. Failures under this context are particularly difficult to mask. They can greatly aggravate the service disruption caused by the upgrade operation itself, and severely affect vendors' reputations. For example, on February 29th, 2012, Azure's service went down after it hit the leap-day bug [28]; in an effort to resolve the issue, developers deployed an upgrade that broke compatibility with a network plugin, causing another three-hour outage.

**Understanding and Detecting Software Upgrade
Failures in Distributed Systems - Zhang et al.,
SOSP 2021**

```
if (message.header.checkpoint_id != self.superblock.working.checkpoint_id() and
    message.header.checkpoint_id !=
        self.superblock.working.vsr_state.checkpoint.parent_checkpoint_id)
{
    // Panic on encountering a prepare which does not match the expected checkpoint
    // id.
    //
    // If this branch is hit, there is a storage determinism problem. At this point
    // in the code it is not possible to distinguish whether the problem is with
    // this replica, the prepare's replica, or both independently.
    log.err(" {}: on_prepare: checkpoint diverged " ++
        "(op={} expect={x:0>32} received={x:0>32} from={})", .{
        self.log_prefix(),
        message.header.op,
        self.superblock.working.checkpoint_id(),
        message.header.checkpoint_id,
        message.header.replica,
    });
    assert(self.backup());
    @panic("checkpoint diverged");
}
```

Another assertion that would catch the bug;
comparing backup's checkpoint ID with primary's

Are we feeling
assertive yet?





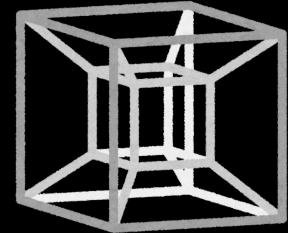
TigerBeetle

TigerSales: Building Trust That Lasts

December 4th, 2025

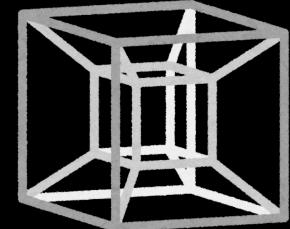
What is trust? (dictionary definition)

Assured reliance on the character, ability, strength, or truth of someone or something.



What is trust? (our own definition)

The undeniable feeling you get when you know someone will raise their hand to take accountability through the good and the bad times. The emotional conviction to know the task at hand will get done well and with unshakeable conviction and character.



What does it feel like to trust?

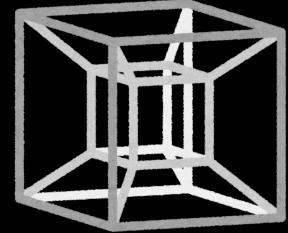
- Your spouse
- Your friends
- Your kids
- Your family
- Your colleagues
- Your manager



At TigerBeetle, we want to build **genuine** trust
that lasts.

Trust is nuanced and deep

The undeniable feeling you get when you know someone will raise their hand to take accountability through the good and the bad times. The emotional conviction to know the task at hand will get done well and with unshakeable character.



Revenue is straightforward, but one-dimensional

The total income produced by a given source.



Most businesses make **revenue** the goal.

They think trust is “**boring**”

But making revenue the goal, **erodes** trust and sets up a transactional relationship.

Building trust ***is*** our goal.

Trust is the input, **Revenue** is the output.

Trust builds revenue, **not** the other way around.

Trust is determined across multiple phases in sales

TRUST IN SALES

- Pre-call research
- Cold outreach engagement
- First call connection
- First meeting follow-up
- Subsequent virtual & in-person meetings
- Team dynamics
- Technology delivery over contract term

What cold outreach from sales looks like today

COLD OUTREACH No prior relationship

Automated

Structured and
templated

Focused on
sender's product

High volume

CTA-centric

“Set it and forget it”

“Use these three options”

“Make sure you list our benefits”

“We’re warming up 20 domains”

“Make sure you get that 15 minute call booked”

What are your thoughts on this cold email?

Technical support metrics that matter pitch

Share

Hit to summarize

From [REDACTED] >
To peter@tigerbeetle.com
Tuesday, November 11 2025 at 6:00 AM PST

Hi Peter,

Running customer success for a high-performance database means your team fields integration questions, API troubleshooting, and technical documentation requests around the clock.

We handle technical support for fintech and database companies - developer onboarding, integration assistance, API documentation questions - so your engineers stay focused on product development instead of support tickets.

Microsoft, Zendesk, and DoorDash use us for their technical customer coordination.

Quick call to discuss developer support scaling for TigerBeetle?

[REDACTED]

Assuming where my pain is vs. knowing

Focused on them vs. me

Assuming I care about these logos

Asking for a meeting on first email

Today's cold email templates are **hurting** your brand
and your chances at building trust.

And it doesn't get much better after the first meeting, either.

"Just following-up on our first meeting.."

"I haven't heard from you so is this not a priority?..."

"Can you let me know which *fits your situation*?:

- A. I offended you in my first meeting
- B. You've been on a vacation island for the past 2 weeks
- C. You're still interested in chatting"

What if you changed how your prospects got to know you?

COLD
No prior relationship

Often automated

Structured and
templated

Focused on
sender's product

High volume

CTA-centric

Highly curated and intentional

Experimental and creative

Focused on trends

High on human connection

Solution-centric

Here's how we'd like to get to know you

COLD
No prior relationship

Highly curated and
intentional

Experimental and
creative

Focused on
universal pains

High on human
connection

Solution-centric



"We're hosting something for the
community..."

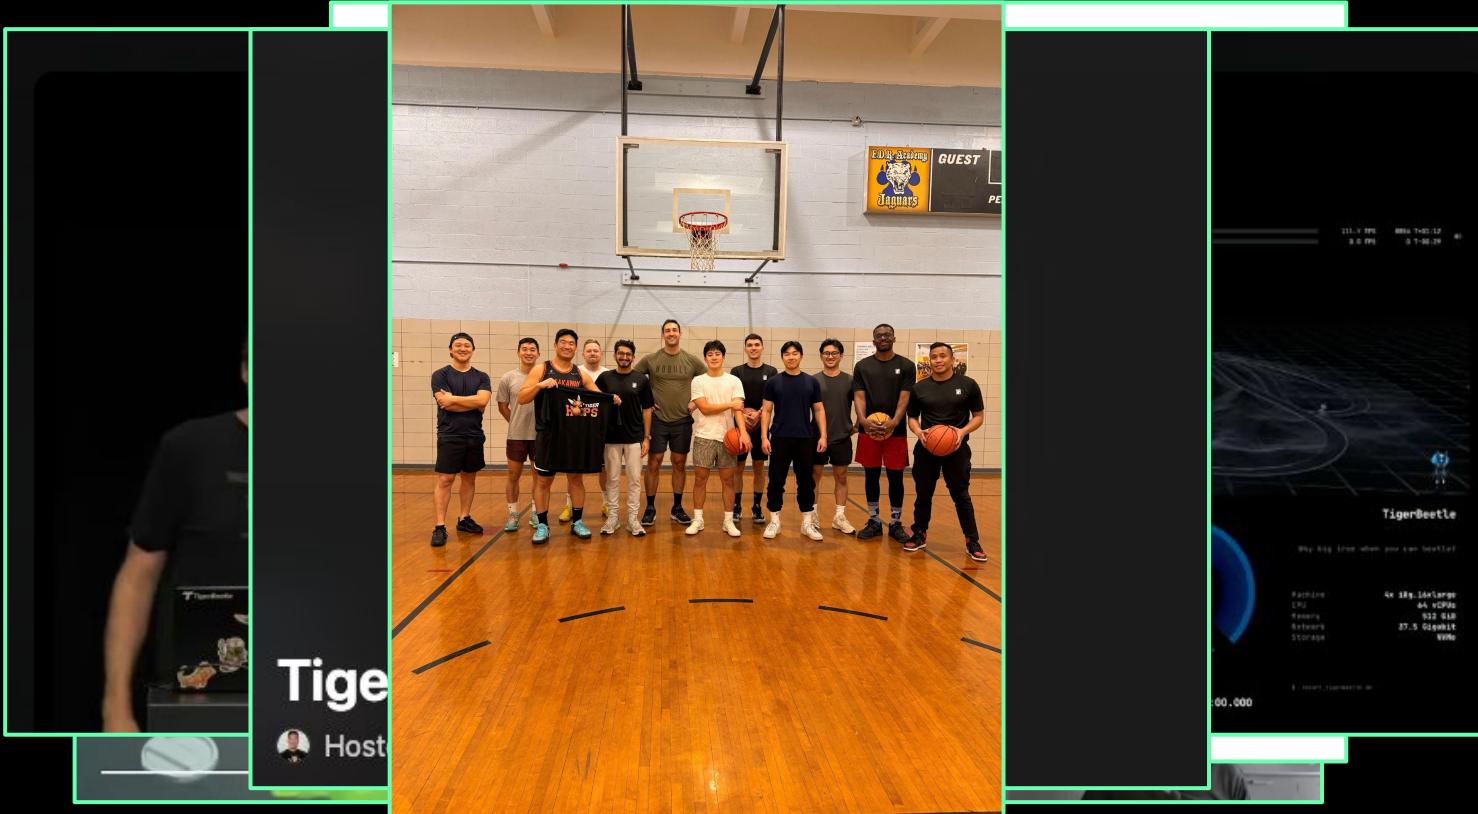
"We created a racer game to demonstrate
TigerBeetle's performance"

"Realtime is only increasing"

"We'd love to meet you..."

"Here's what we believe is the
path..."

BTHF: Be There Have Fun!



What you won't see us do

- X Automated cold sequences
- X Multiple follow-ups after the first call



What we *will* do

Continue to bring the community together

Create authentic 1:1 relationships

BTHF!

TigerBeetle



Thank you!

APPENDIX

Now think about the technology partners and vendors you work with...do you **trust** them?

We'd like to share a thing or two about how we think
about **building** it.

Trust is also **hard** to measure, takes time, and is subjective.

Trust is determined across multiple phases in sales

TRUST

- Pre-call research
- Cold outreach engagement
- First call connection
- First meeting follow-up
- Subsequent virtual & in-person meetings
- Team dynamics
- Technology delivery over contract term

REVENUE

- Post contract transaction

What you won't see us do

- X Automated cold sequences
- X Multiple follow-ups after the first call
- X Sticking around when relationships lack mutual respect
- X Compromising our values when faced with a shiny logo



What we *will* do

Continue to bring the community together

Create authentic 1:1 relationships

Help support the transition to realtime

BTHF!

