



T TigerBeetle

1000X WORLD TOUR

13 CITIES | 6 DAYS



TigerFans: High-Performance Ticketing with TigerBeetle

From Double-Entry Accounting to 977 Tickets per Second

The screenshot shows a dark-themed website for "TigerFans 2025". At the top left is the "TigerFans" logo with a star icon. At the top right are links for "TigerBench" and "GitHub". The main title "TigerFans 2025" is prominently displayed in the center. Below it, a subtitle reads "A fictional conference for people who love fast, correct systems." and the date "Dec 3–4, 2025". Two buttons at the bottom are labeled "Order now" and "Why attend?". To the right, there's a large callout box containing technical details: "powered by TigerBeetle — 1000x FASTER OLTP", "reservations store Redis", and "orders database PostgreSQL". It also mentions "tigerbeetle cluster size 1 (single-replica)". A note at the bottom states: "All transactions in the demo are mock-paid, inventory tracked with TigerBeetle."



It began with a tweet.

"How would you sell Oasis-scale tickets without overselling?"

Answer from Joran: Too easy: TigerBeetle.

I wanted to know:

- **HOW** do you actually model ticketing as **double-entry accounting**?
- What does the **real code** look like?
- How do you handle **timeouts, webhooks, idempotency**?

So I decided to build the whole thing.



I didn't want a toy example.

I wanted a **realistic ticketing flow**

- Checkout with multiple ticket classes & goodies
- Pending reservations with timeouts
- Webhook-based payment confirmation
- Strict no overselling guarantee

Stack:

- FastAPI (Python, async)
- TigerBeetle (accounting)
- SQLite at first
- MockPay (fake payment provider)



TigerBeetle gives us **accounts, transfers, debits, credits.**

We turn ticketing into an **accounting problem:**

- Every reservation is a **transfer**
- Every ticket class is a set of **accounts**
- **Debits = Credits => system never goes out of balance**

Double-entry accounting gives us:

- Built-in **error detection**
- **Auditability for free**
- And with TB: **durability + performance in the same model**

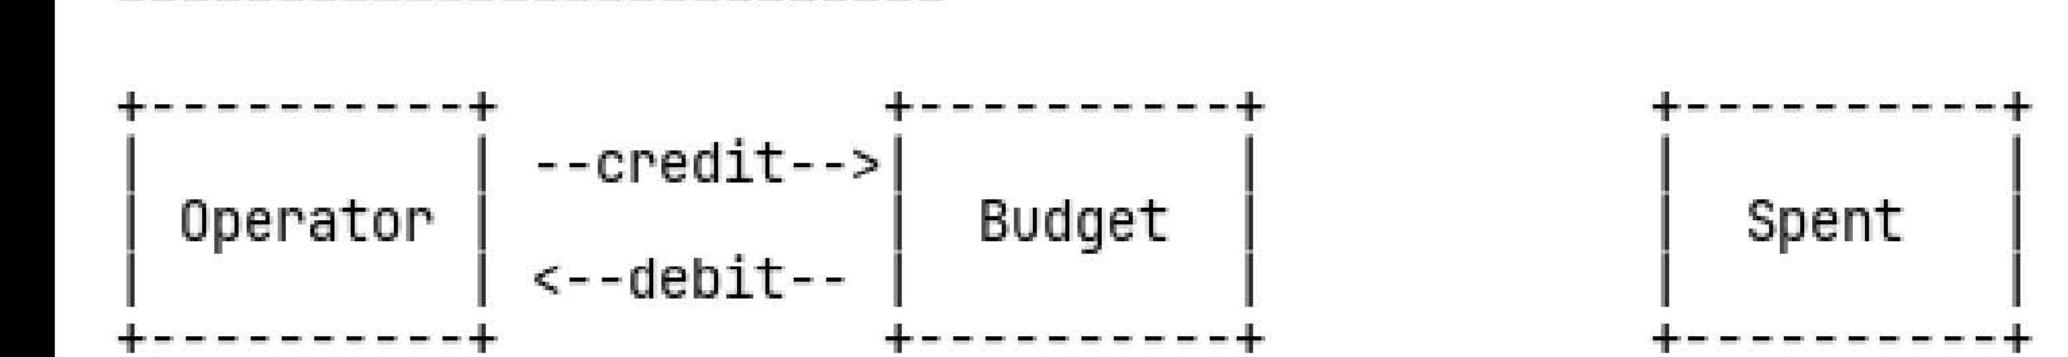


Three Accounts per Ticket Class

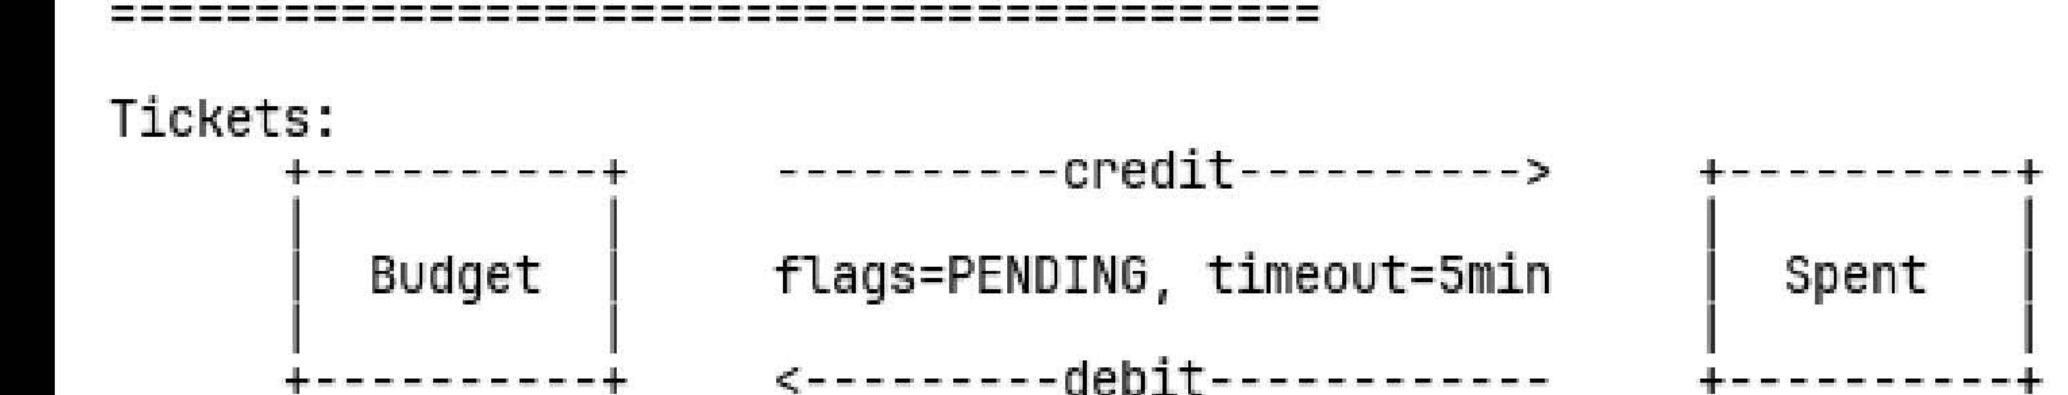
For each ticket class:

- **Operator:**
Holds all inventory
- **Budget:**
What we're allowed to sell
- **Spent:**
Tickets actually reserved/used

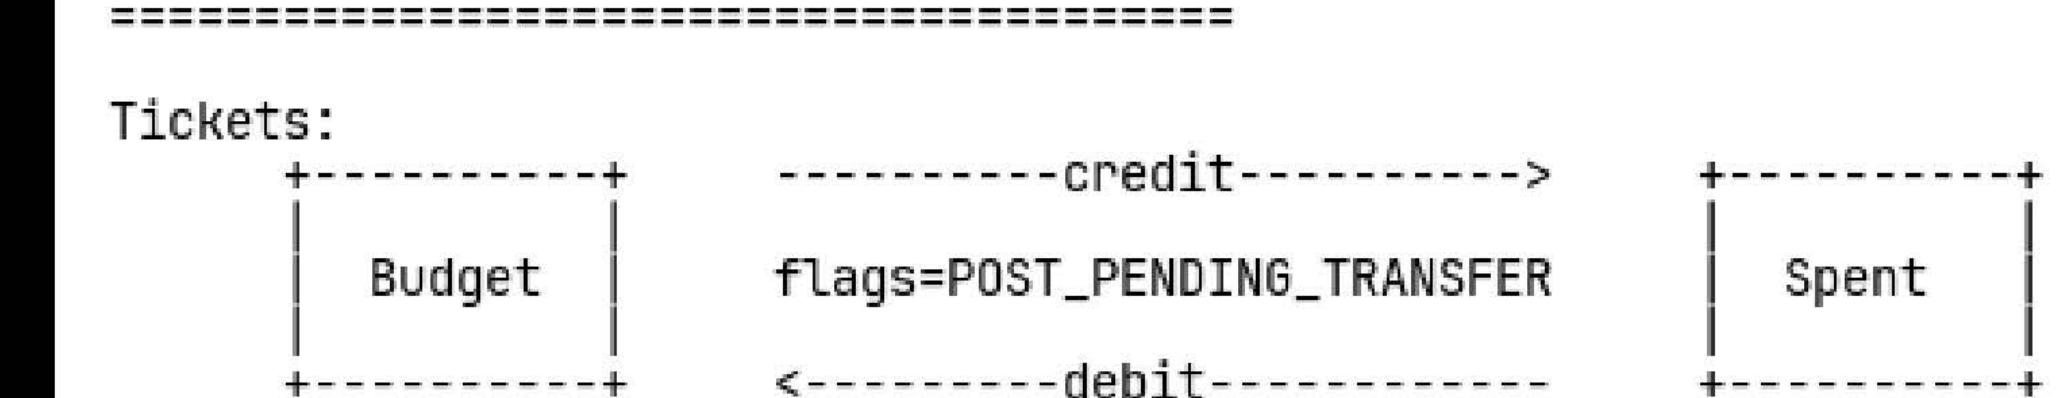
INIT FLOW: Fund the budget



HOLDING FLOW: Limited Hold on the resource



COMMIT FLOW: Post the pending transfers



Flow:

1. Fund Budget from Operator
2. Move from Budget Spent on reservation
3. Post or void based on payment outcome



Pending Transfers & Timeouts

Checkout = create pending transfer:

- Budget -> Spent
- 5-minute timeout
- DEBITS MUST NOT EXCEED CREDITS flag on accounts

Payment succeeds:

- Post the pending transfer => ticket is locked in

Payment fails / user disappears:

- Void the transfer => it just vanishes

No cron jobs. No cleanup races.

Correctness is enforced by TigerBeetle's invariants.

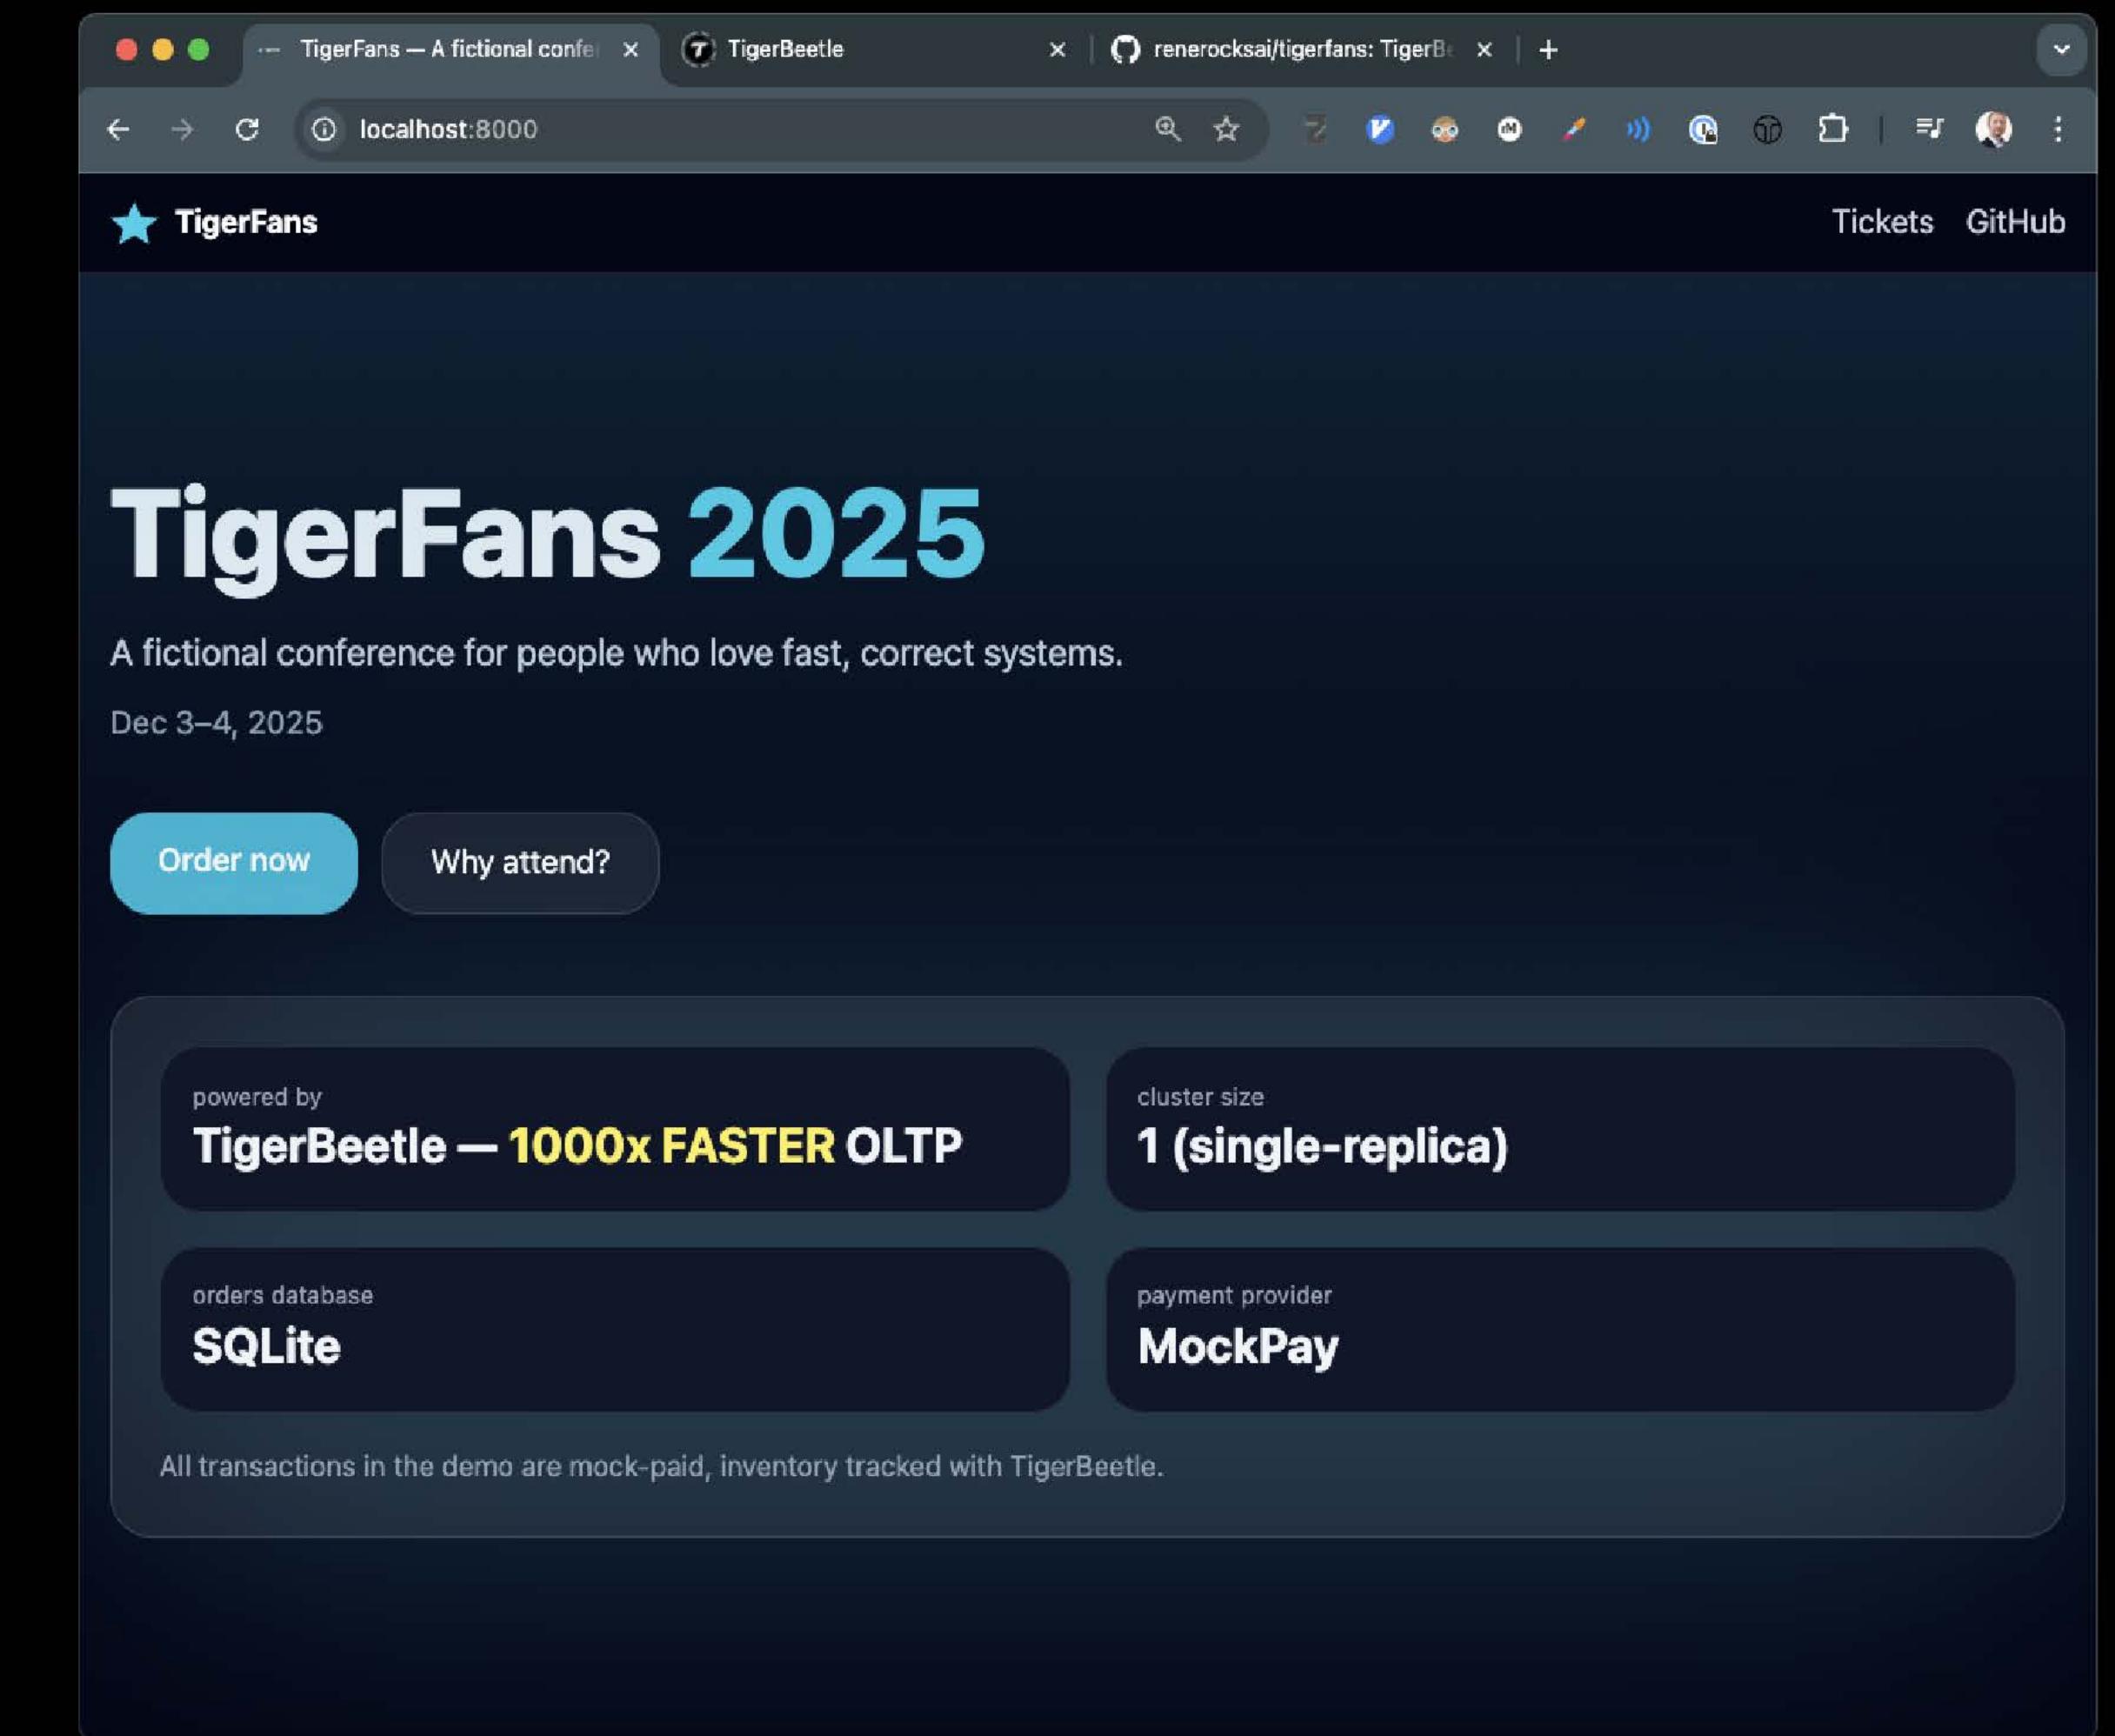


Version 1:

- FastAPI (1 worker)
- SQLite
- TigerBeetle (dev mode)
- MockPay

Everything worked:

- 2-phase checkout
- Timeouts
- Webhooks
- UI + Admin view





First measurement:

- **~115 tickets/sec**

Oasis baseline:

- **~1.4M tickets / 6 hours**
- **≈ 65 tickets/sec**

So we're at about **O(1.7 Oasis)**.

I send this to Joran...



Baseline Performance

First measurement:

- ~115 tickets/sec

Oasis baseline:

- ~1.4M tickets / 6 hours
- ≈ 65 tickets/sec

So we're at about O(1.7 Oasis).

I send this to Joran...

...he replies:

"I was surprised the TPS is so low.

It should be ~10k.

Where Is the Time Going?

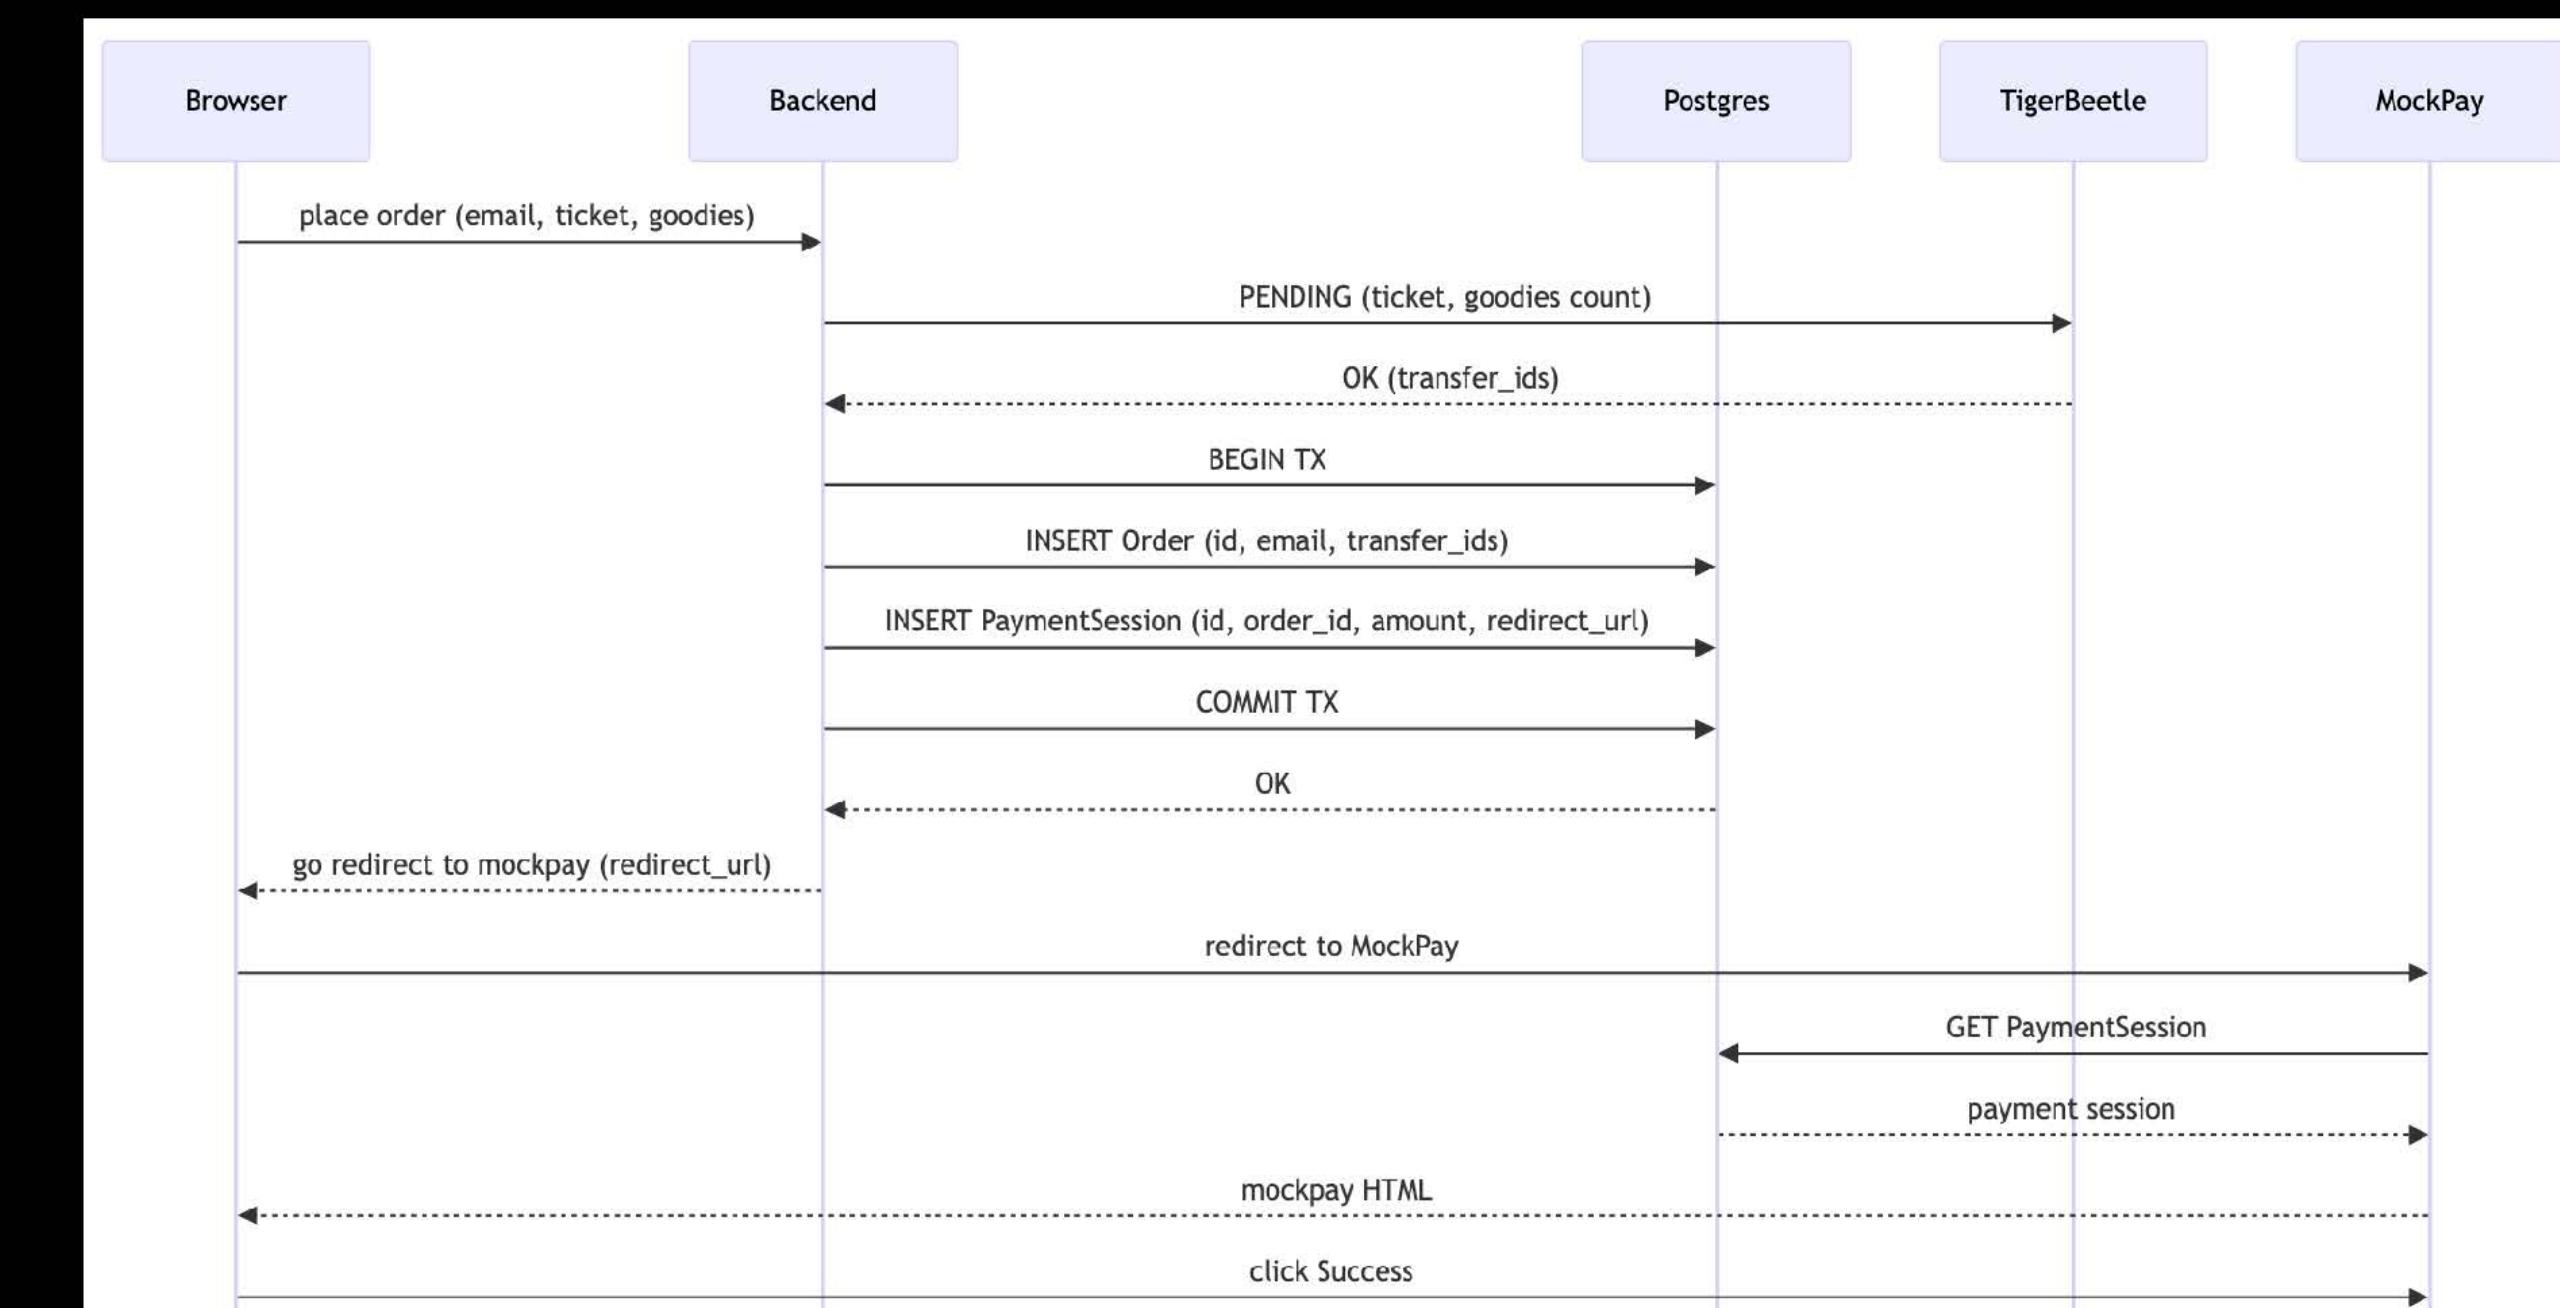


Checkout flow:

1. Browser -> Server: *place order*
2. Server -> TB: PENDING transfers
3. Server -> DB: sessions + orders
4. Server -> Browser: redirect

Webhook flow:

1. MockPay -> Server
2. Server -> DB: idempotency checks
3. Server -> TB: POST transfers
4. Server -> DB: final order write





Bottleneck #1: PostgreSQL Everywhere

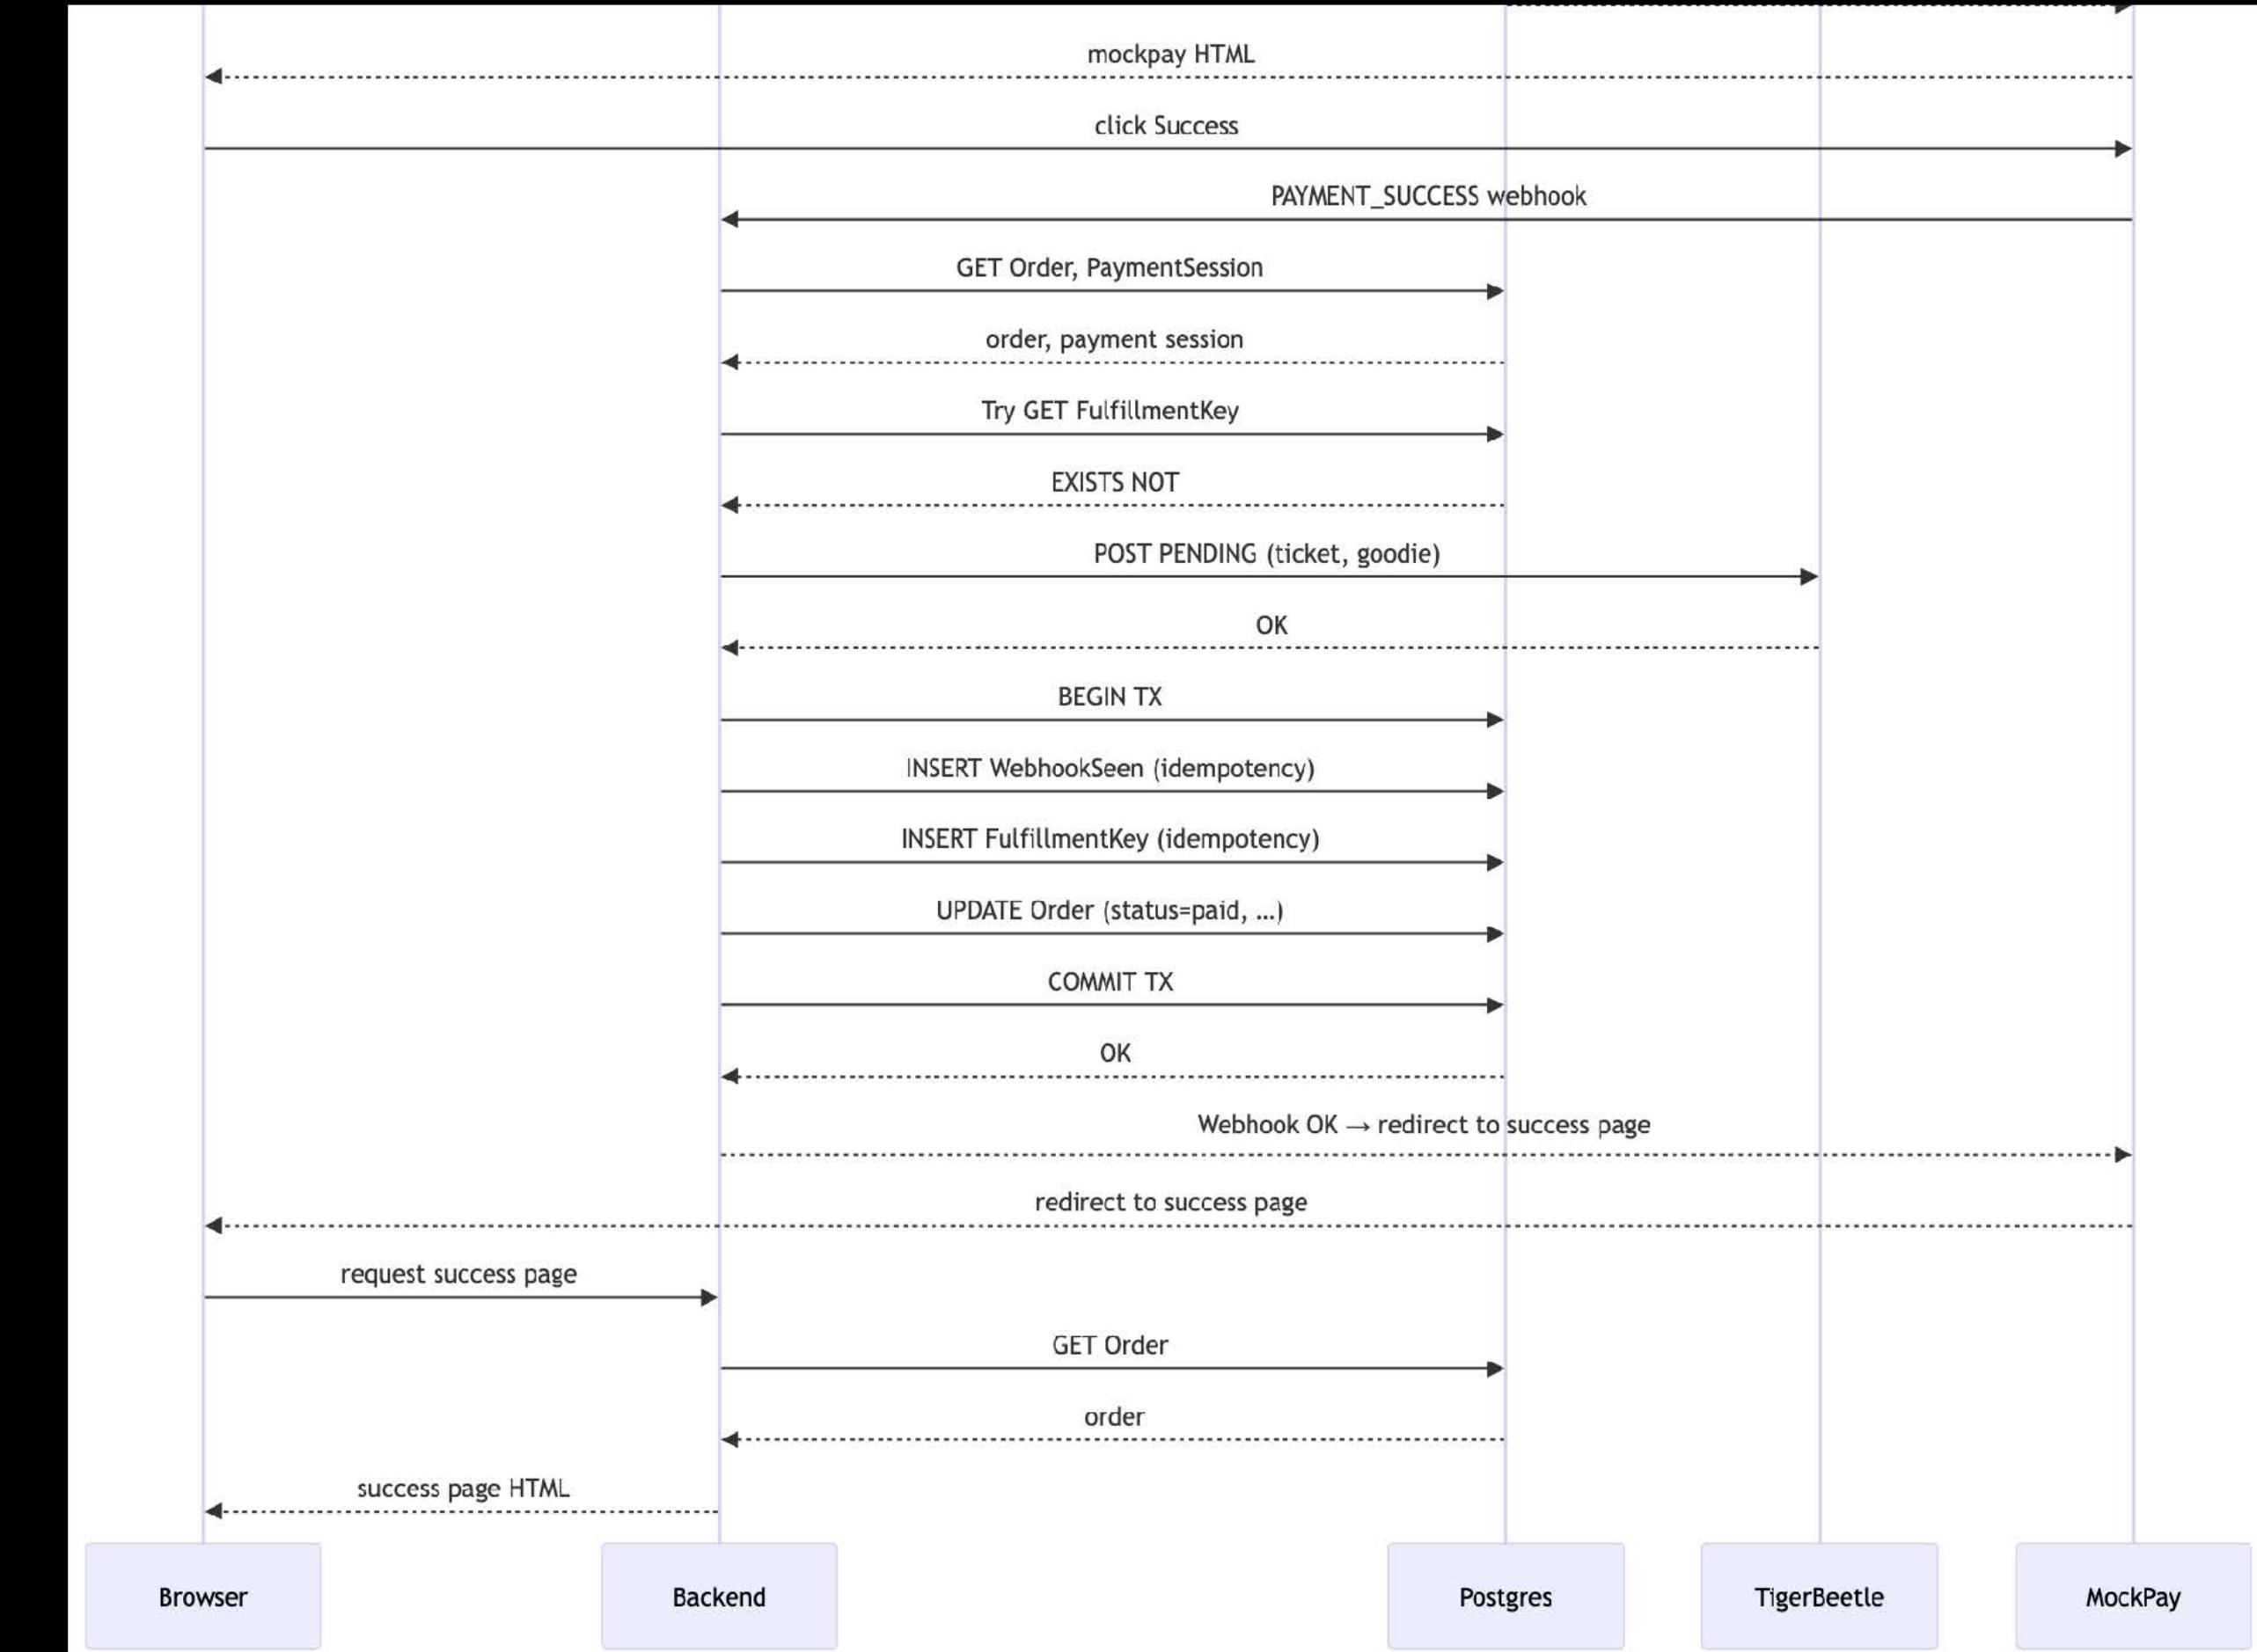
With PostgreSQL in the hot path:

- Reservations: ~150 ops/s
- Webhooks: ~130 ops/s

Every request:

- 2–4 round-trips to PG
- Multiple fsyncs
- Contention on the same tables

The database, not TigerBeetle, is the bottleneck.





Redis Experiment: All-In Memory

I swap PostgreSQL -> Redis for **everything**:

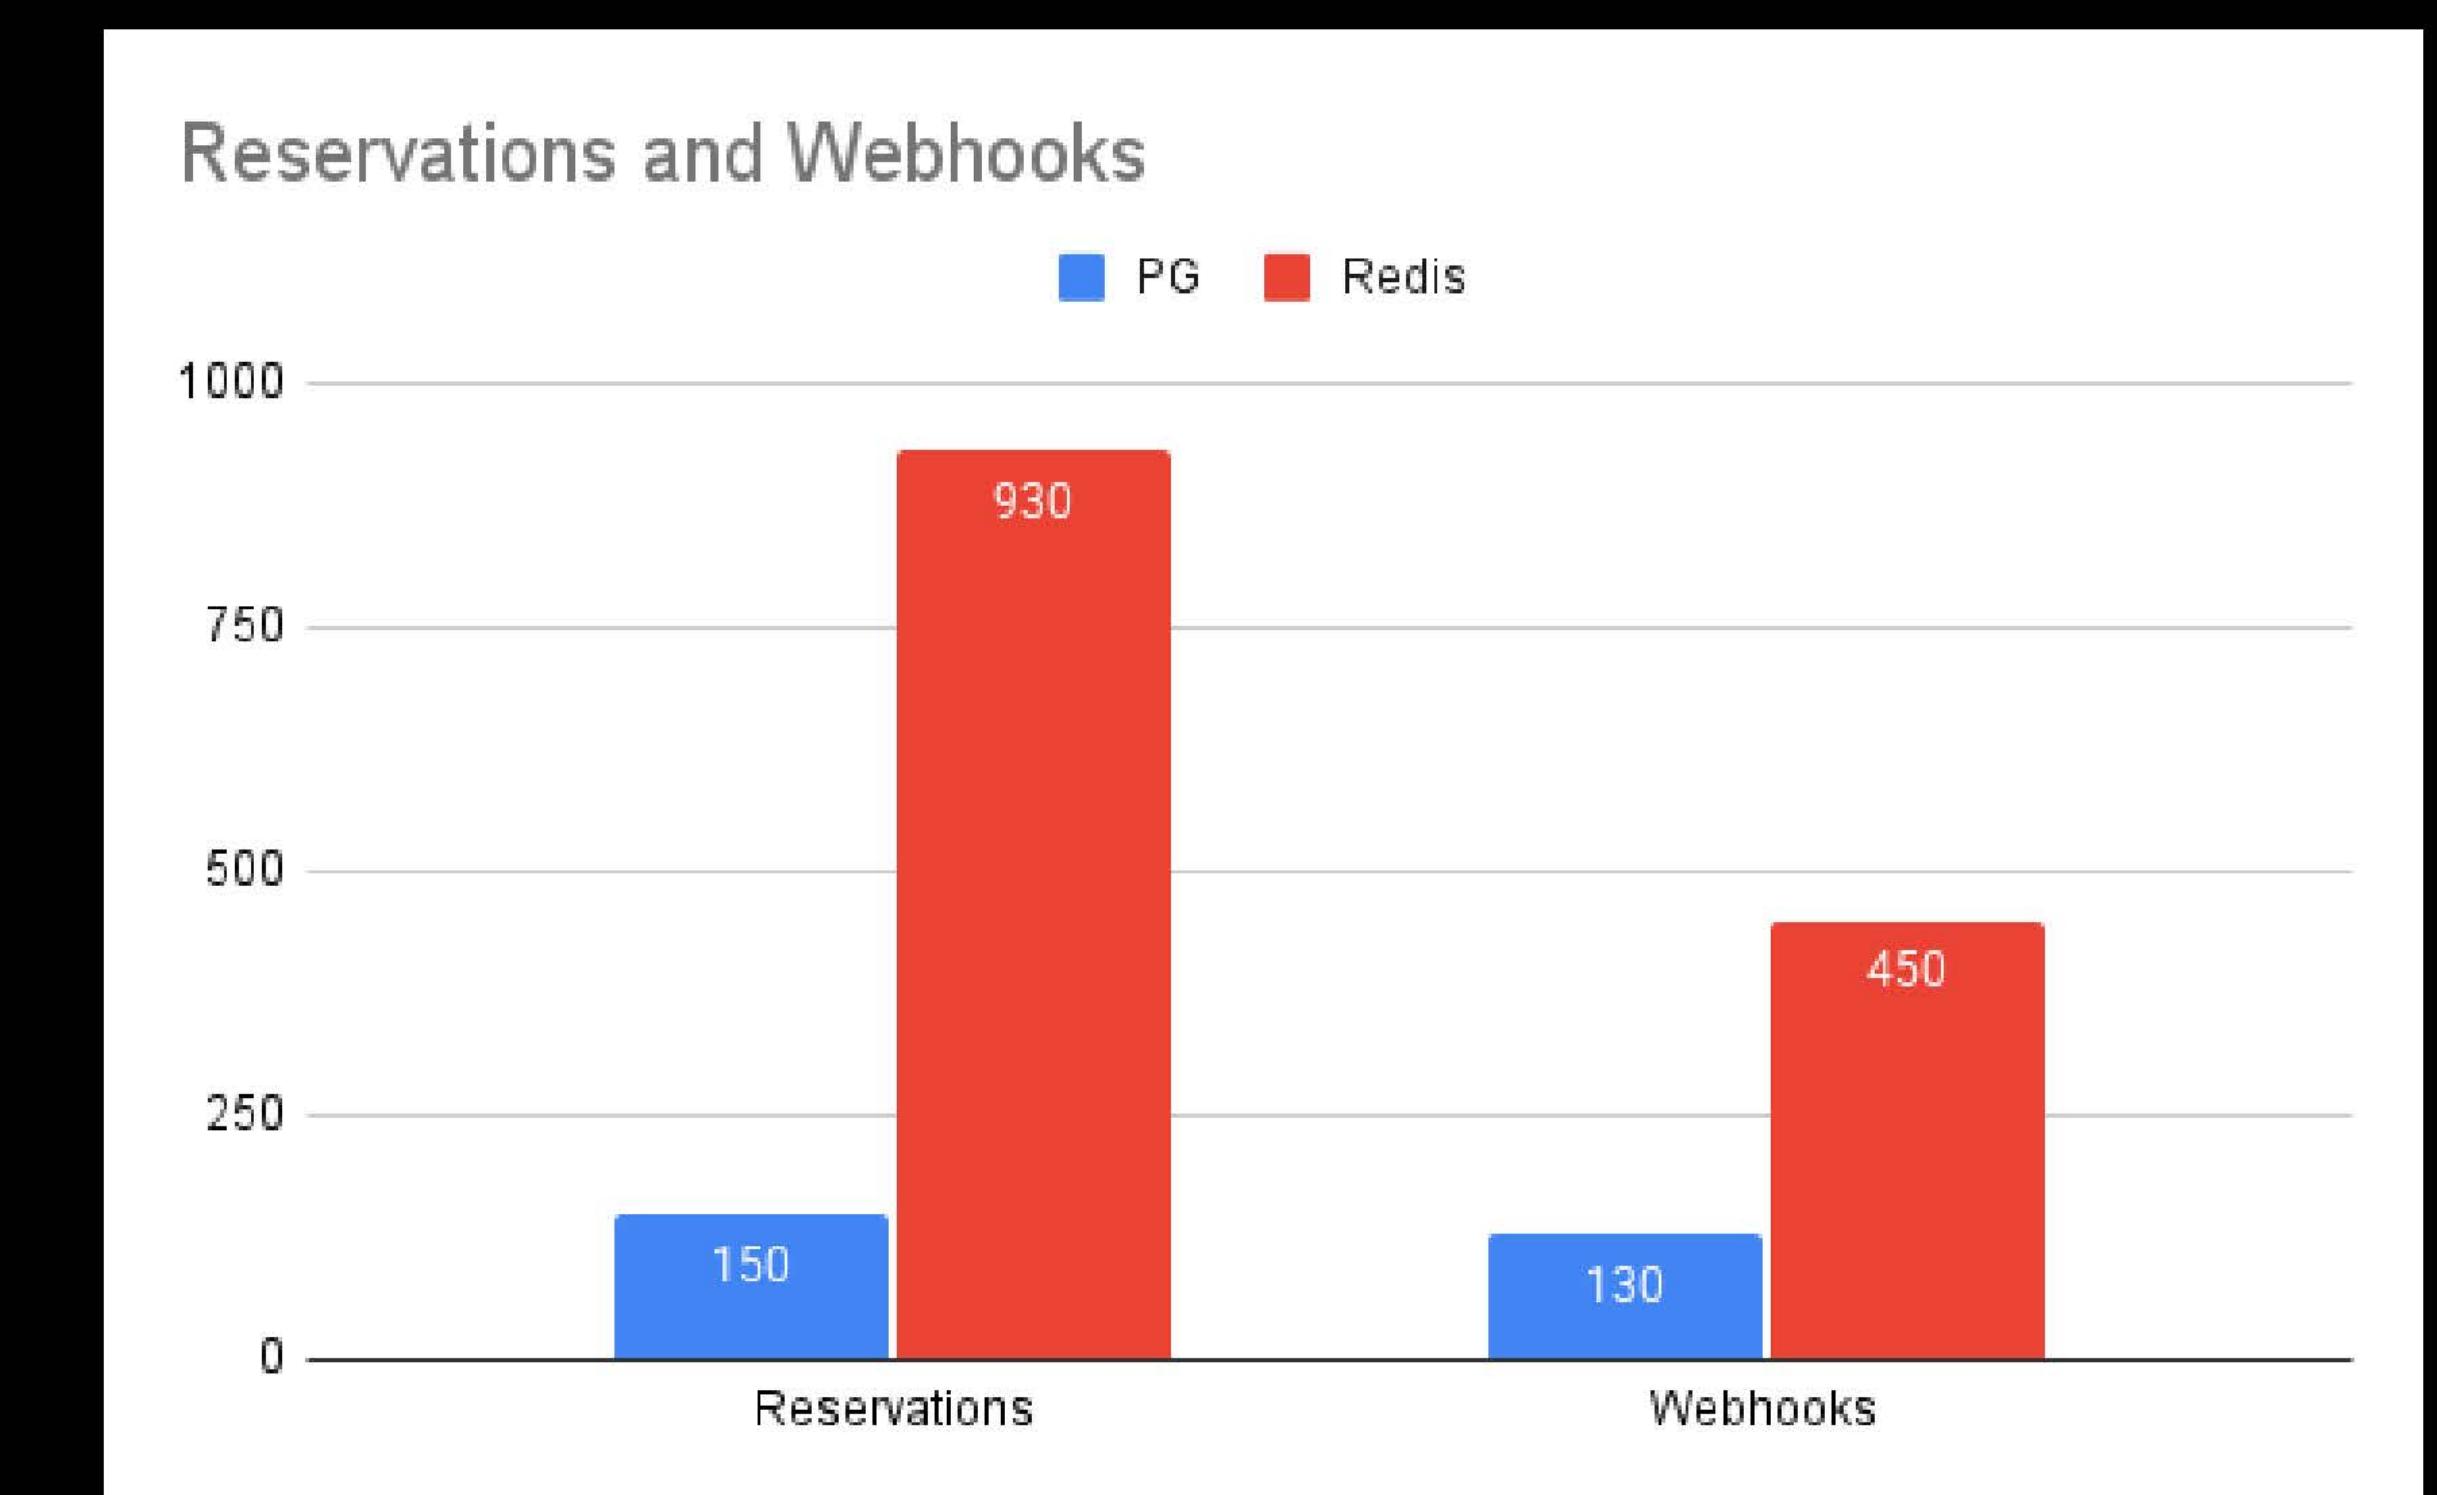
- Sessions
- Idempotency
- Orders

Results:

- Reservations: ~930 ops/s
- Webhooks: ~450 ops/s

Great... except Redis **everysec can lose 1s of orders** on crash.

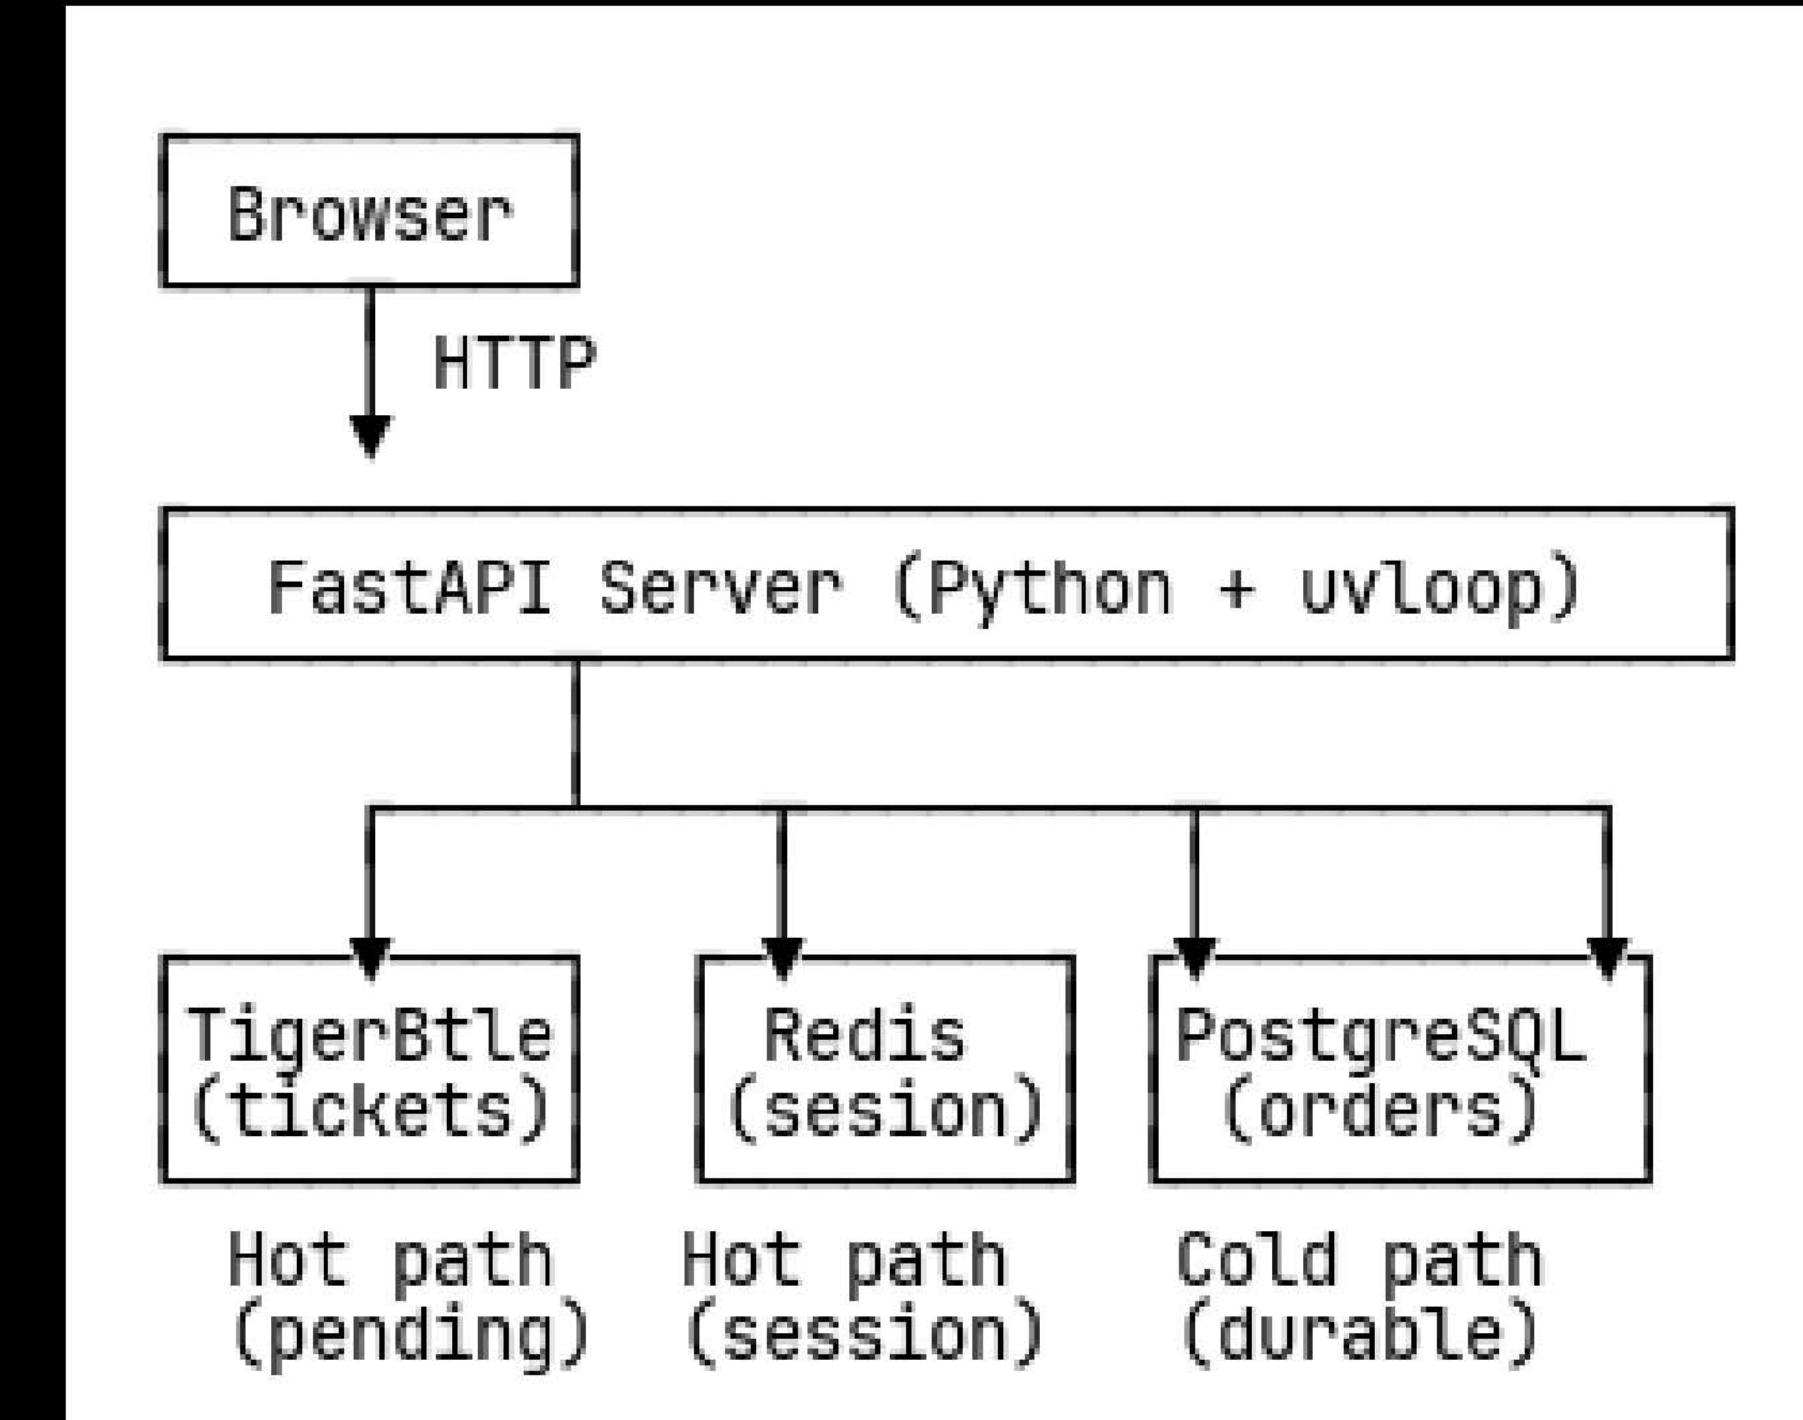
That's not acceptable.





Rafael's suggestion:

- Use **Redis** for **hot data**:
 - Payment sessions
 - Idempotency keys
- Use **PostgreSQL** only for **cold data**
 - Completed orders



TigerBeetle:

- Always handles **ticket accounting**
- Always fully durable



Three Configuration Levels

Level 1 – PG Only

- PG for sessions, accounting, orders

Level 2 – PG + Redis

- Redis: sessions
- PG: accounting + orders

Level 3 – TB + Redis

- Redis: sessions
- TigerBeetle: accounting
- PG: orders





Interface Impedance: Requests vs Batches

TigerBeetle loves big batches:

- Up to 8,190 transfers per call
- One network round-trip amortized

FastAPI loves per-request awaits:

- Each request calls `create_transfers()` once
- No chance to “wait a bit and batch”

Instrumentation showed:

- Batch size ≈ 1
- We were flying a 747 with 1 passenger per flight.

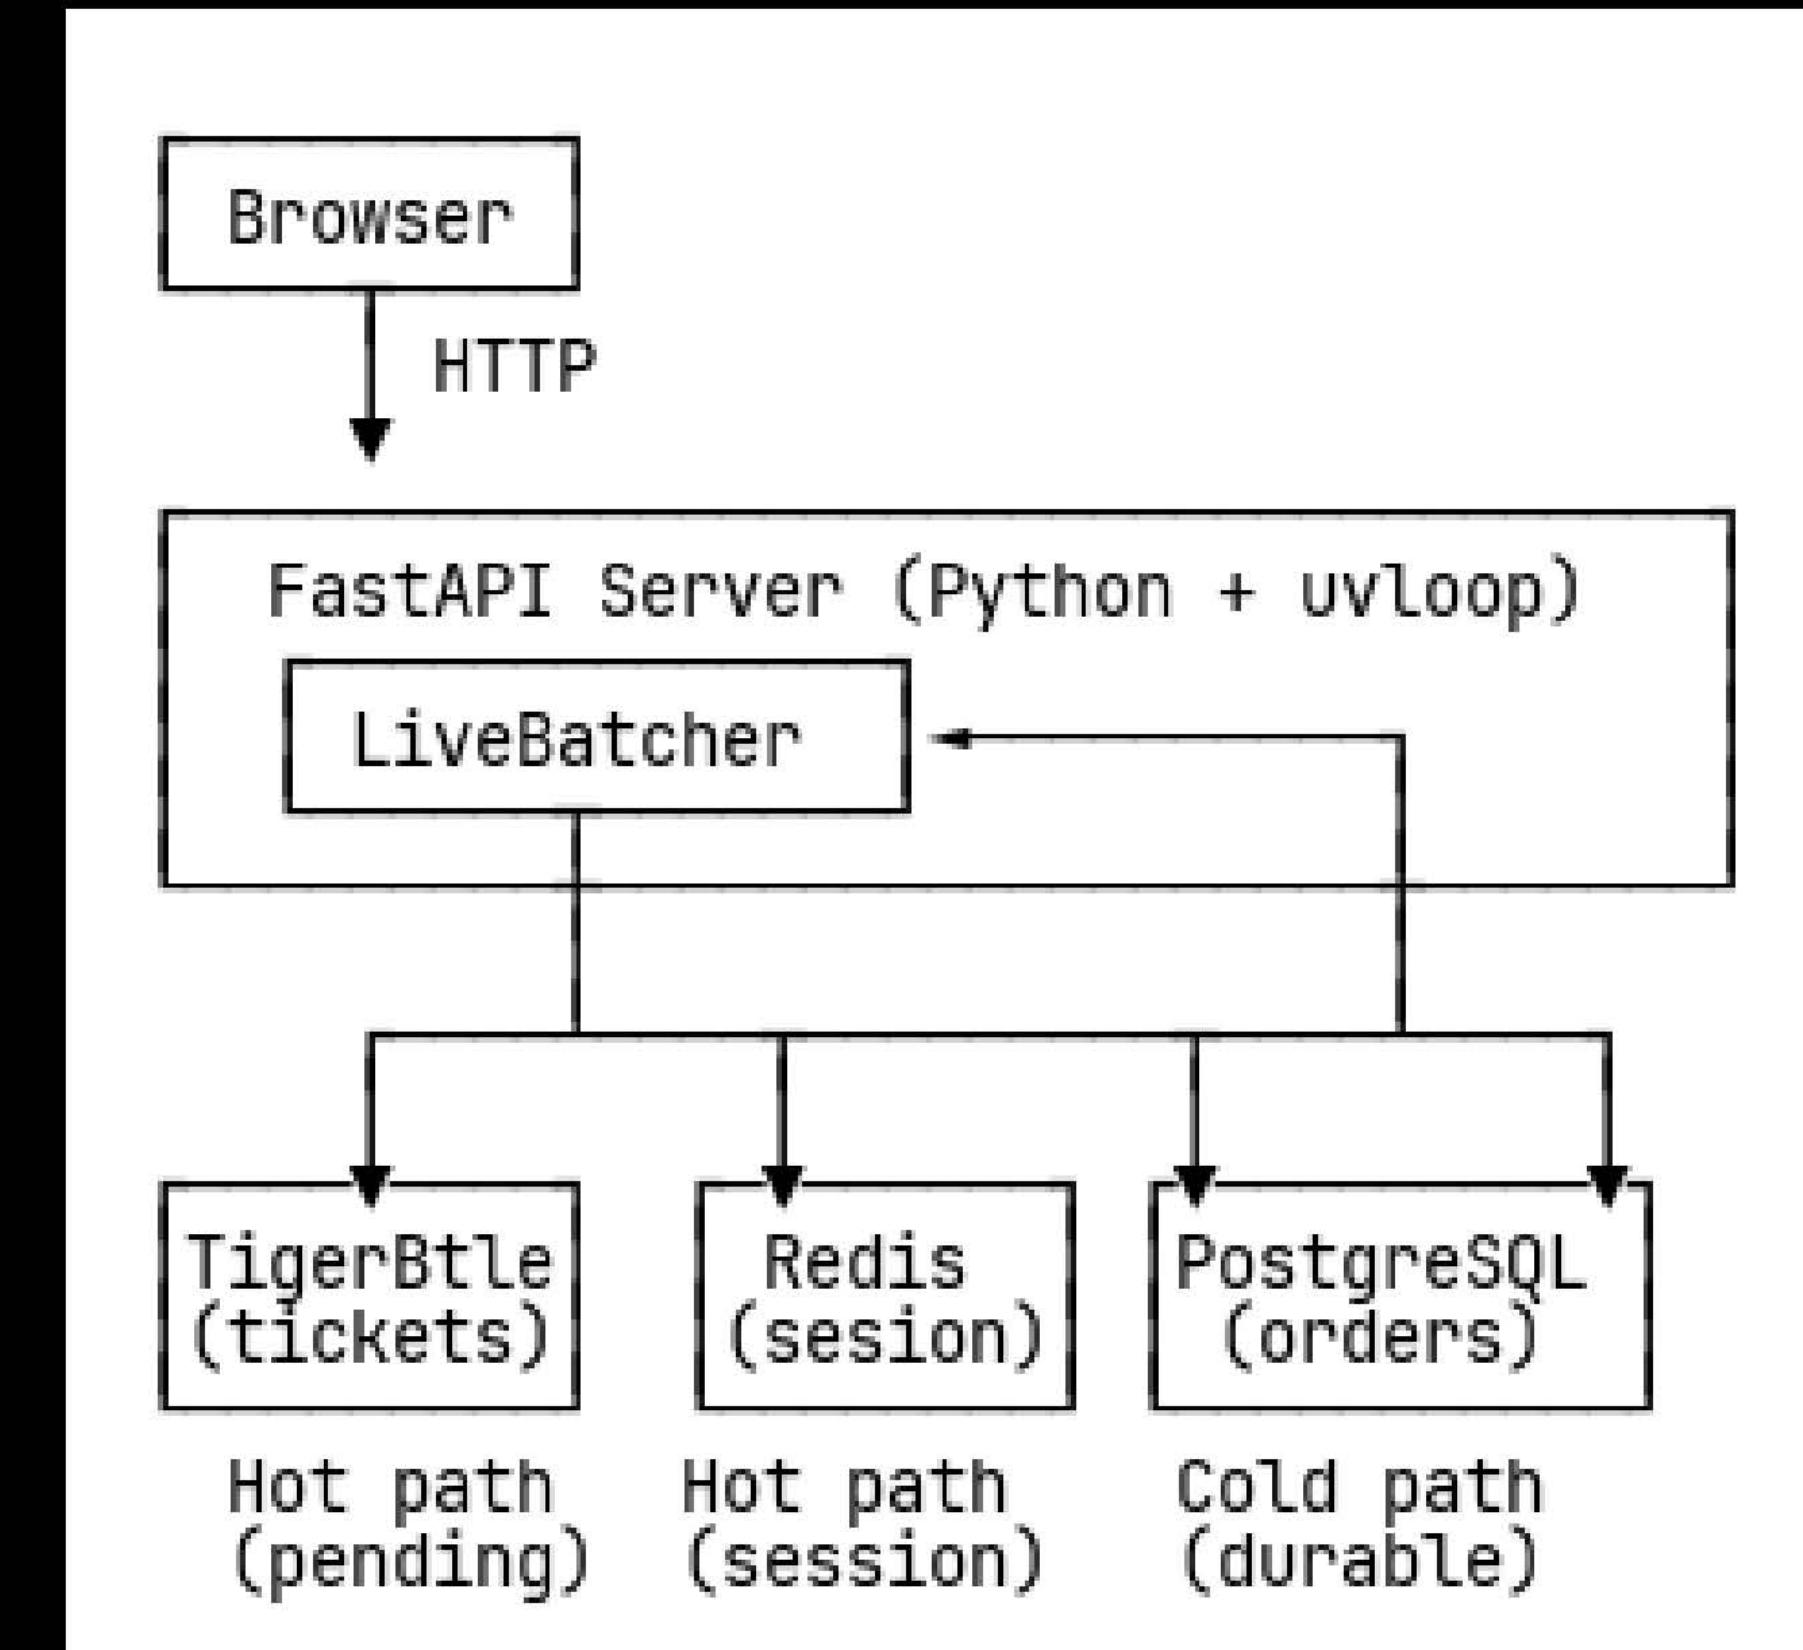


LiveBatcher:

- Sits between app and TB client
- **Collects concurrent requests**
- **Packs them into batches while previous batch is in flight**
- Maps TB errors back to each caller

Behavior:

- Under load: **batch sizes 5–6 transfers**
- Low load: small batches, low latency





Batching Results

With hot/cold path + LiveBatcher:

- Reservations: **977 ops/s**
- Average TB **batch size: 5–6**

Compared to baseline:

- PG-only: ~150 ops/s
- TB+Redis (before batching): ~900 ops/s

LiveBatcher gave us the final 8% – the earlier wins came from architecture, not micro-optimizations.



The Single-Worker Paradox

1 worker

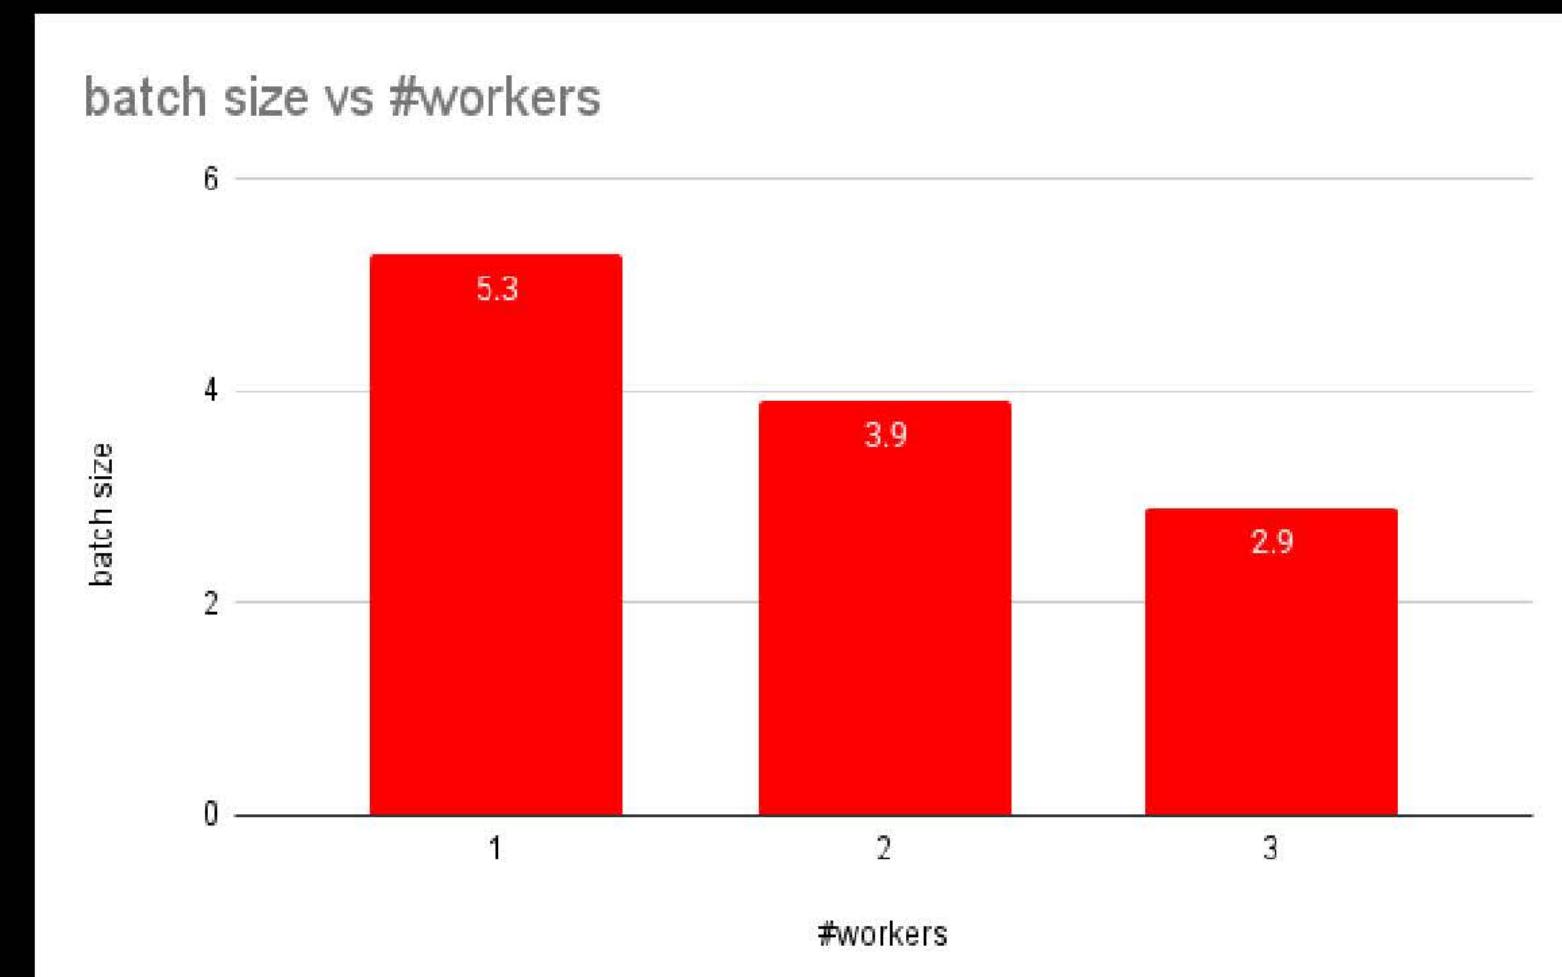
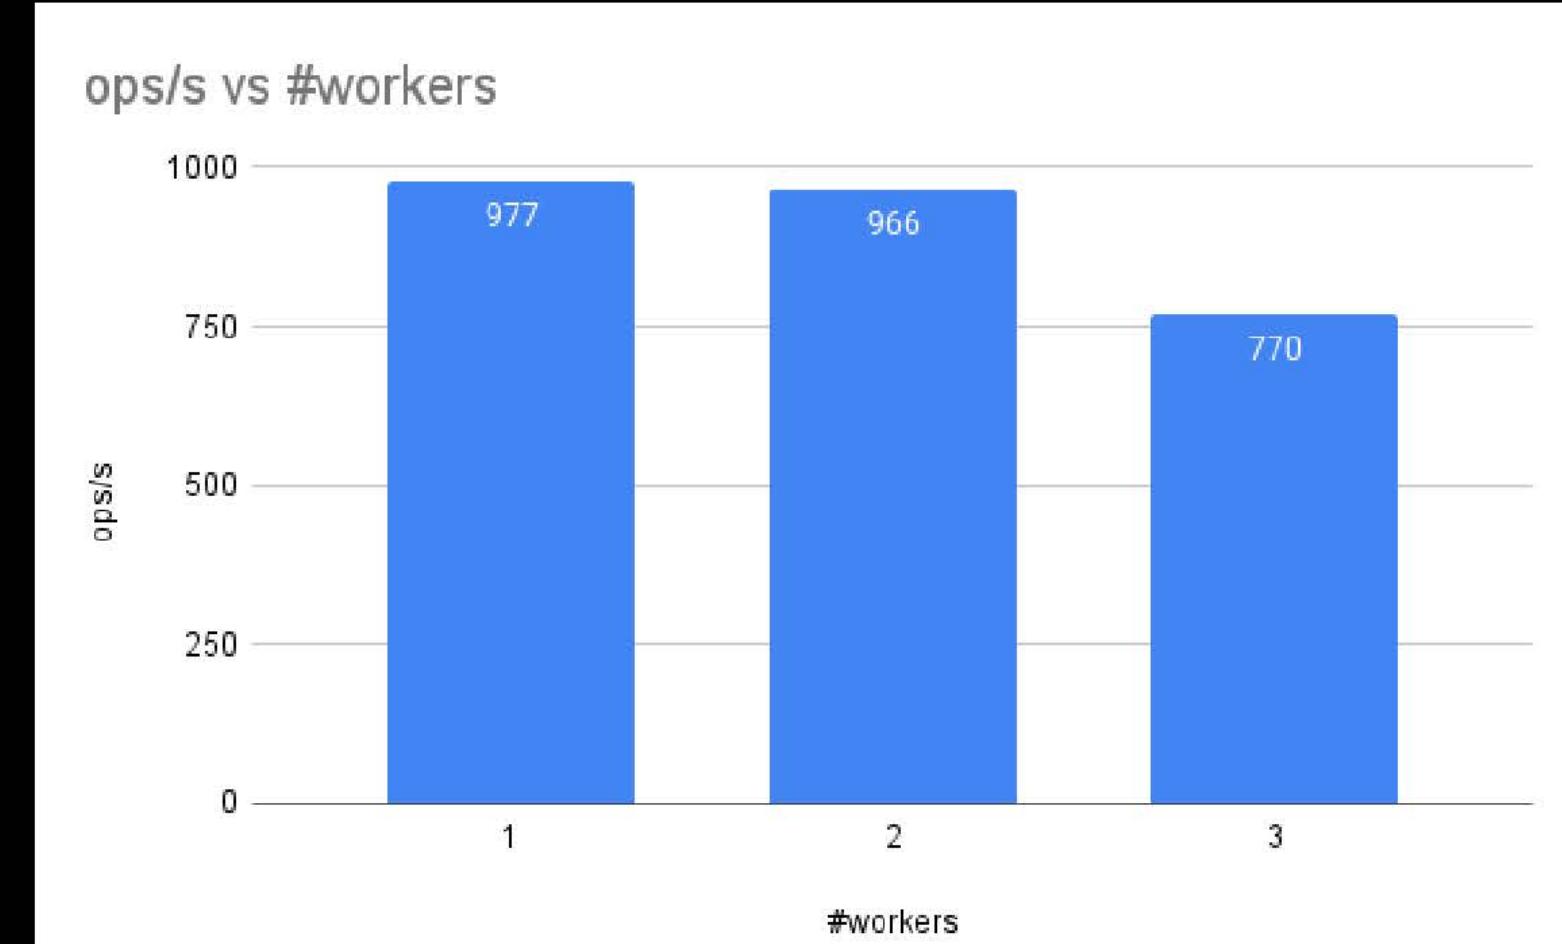
- 977 ops/s
- Avg batch size: 5.3

2 workers

- 966 ops/s
- Avg batch size: 3.9

3 workers

- 770 ops/s
- Avg batch size: 2.9



More workers -> smaller batches -> lower throughput.



End result (Python, 1 worker):

- **977 reservations / second**
- $\approx 15\times$ Oasis baseline
- TB batch size: **5–6 transfers**

The real limit now isn't TigerBeetle
— it's Python's event loop overhead.



In checkout:

- We sped up **accounting** (PG TB) by $\approx 3\times$
- Moved sessions to Redis
- But every request still spends:
 - $\sim 5\text{ms}$ in Python/FastAPI
 - Routing, JSON, business logic

Amdahl's Law says:

- That **serial 5ms chunk becomes a hard ceiling**
- Even with “infinite-speed” TB, Python caps us



What Amdahl Suggests Next

If we keep the same architecture:

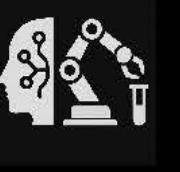
- Same double-entry model
- Same hot/cold split
- Same LiveBatcher idea

...but move to:

- Go, Zig, ...

Then:

- Python's 5ms -> maybe 0.1 – 0.5ms
- Batch sizes can grow
- Amdahl's Law says:
- 10–50× more throughput is realistic



CHALLENGE:

- Same patterns: <-- or do it your way :-)
 - Double-entry resource modeling
 - Hot/cold path
 - Auto-batching
 - Single-worker friendly design
- Any language, any stack <-- ... any platform? Raspberry Pi?
- Benchmark it, share results <-- please :-)

Resources:

- tigerfans.io (demo + TigerBench)
- github.com/renerocksai/tigerfans



1. Ticket Booking

- Start a checkout
- See pending reservation appear

2. MockPay

- Complete payment
- Webhook posts TB transfers

3. Admin Live View

- Orders, reservations, timeouts

4. TigerBench

- Compare PG-only vs PG+Redis vs TB+Redis
- Ops vs. Transactions



TigerFans: High-Performance Ticketing with TigerBeetle

From Double-Entry Accounting to 977 Tickets per Second

THANK YOU for your attention!



[Backup] Amdahl's Law: Component Breakdown

Baseline checkout (Level 1 – PG Only):

- Total time ≈ 15.4 ms
- PG accounting: **14.03 ms** ($\sim 91\%$)
- PG sessions: **1.36 ms** ($\sim 9\%$)

Optimized checkout (Level 3 – TB + Redis):

- Total time ≈ 5.63 ms
- TB accounting: **4.58 ms**
- Redis sessions: **1.04 ms**

Amdahl's Law:

$$\text{Speedup_overall} = 1 / ((1 - P) + P / S)$$

Here:

- Accounting $P \approx 0.91$, $S \approx 3.0$
- Predicted speedup $\approx 2.73\times$
- Measured component speedup $\approx 2.74\times$



[Backup] Amdahl & the Batching Ceiling

For batching, Amdahl applies again:

- TigerBeetle batch time ≈ 3 ms
- Python per-request serial work ≈ 3 ms
- Routing, parsing, validation
- Business logic
- Preparing TB transfers

During one TB batch (~ 3 ms):

- Python can push ≈ 1 new request to the batcher

So batch size naturally stabilizes around:

- 2–6 transfers depending on concurrency load

In Go/Zig:

- Serial overhead ≈ 0.1 ms
- Same architecture => batches of 30–50+